

## **\*\* Charts and Tensorboard plots are visible in the pdf version submitted \*\***

### **Assignment 5 - RNN, LSTM and GRU models**

In [ ]:

In [256]: `from sklearn.ensemble import VotingClassifier`

In [162]: `import os`

In [ ]:

In [ ]:

In [65]: `import keras  
from keras.optimizers import SGD`

In [66]: `import re  
from keras.models import Sequential, load_model  
from keras.layers import Dense, LSTM, Embedding, Dropout  
from keras.preprocessing.text import Tokenizer  
from keras.preprocessing.sequence import pad_sequences`

In [ ]:

In [67]: `import tensorflow as tf  
from tensorflow.keras import layers`

In [68]: `import datetime`

In [69]: `from sklearn.model_selection import train_test_split`

In [ ]:

```
In [70]: from keras import optimizers
```

```
In [71]: import numpy as np
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

### Task 1 - Load and preprocess the data

```
In [72]: import pandas as pd

# Read train data
df = pd.read_csv('data/sst_train.txt', sep='\t', header=None, names=['truth', 'text'])
```

```
In [ ]:
```

```
In [73]: # Read dev data
dev = pd.read_csv('data/sst_dev.txt', sep='\t', header=None, names=['truth', 'text'])
```

```
In [ ]:
```

```
In [74]: # Read test data
test = pd.read_csv('data/sst_test.txt', sep='\t', header=None, names=['truth', 'text'])
```

```
In [ ]:
```

```
In [ ]:
```

```
In [75]: df['truth'] = df['truth'].str.replace('__label__', '')
```

In [ ]:

In [76]: `dev['truth'] = dev['truth'].str.replace('__label__', '')`

In [ ]:

In [77]: `test['truth'] = test['truth'].str.replace('__label__', '')`

In [ ]:

In [78]: `df['truth'] = df['truth'].astype(int).astype('category')`

In [ ]:

In [79]: `dev['truth'] = dev['truth'].astype(int).astype('category')`

In [ ]:

In [80]: `test['truth'] = test['truth'].astype(int).astype('category')`

In [ ]:

In [81]: `dev.head()`

Out[81]:

	truth	text
0	4	It 's a lovely film with lovely performances b...
1	3	No one goes unindicted here , which is probabl...
2	4	And if you 're not nearly moved to tears by a ...
3	5	A warm , funny , engaging film .
4	5	Uses sharp humor and insight into human nature...

In [ ]:

In [82]: `test.head()`

Out[82]:

	truth	text
0	3	Effective but too-tepid biopic
1	4	If you sometimes like to go to the movies to h...
2	5	Emerges as something rare , an issue movie tha...
3	3	The film provides some great insight into the ...
4	5	Offers that rare combination of entertainment ...

In [ ]:

In [ ]:

In [83]: `df.head()`

Out[83]:

	truth	text
0	4	The Rock is destined to be the 21st Century 's...
1	5	The gorgeously elaborate continuation of `` Th...
2	4	Singer/composer Bryan Adams contributes a slew...
3	3	You 'd think by now America would have had eno...
4	4	Yet the act is still charming here .

In [ ]:

In [84]: `# count training samples`  
`df.shape[0]`

Out[84]: 8544

In [ ]:


In [ ]:

In [ ]:

In [85]: `import matplotlib.pyplot as plt`

## Dev data Visualizations

```
In [86]: ax = dev['truth'].value_counts(sort=False).plot(kind='barh')
ax.set_xlabel("Number of Samples in dev Set")
ax.set_ylabel("Label")
plt.show()
```



In [ ]:

In [87]: `data_dev = dev['text']`In [88]: `dev_X = list()`

```
In [89]: for i in range(len(dev)):  
         integer_encoded = [ord(char) for char in data_dev[i]]  
         dev_X.append(integer_encoded)
```

```
In [ ]:
```

```
In [90]: dev_X_test = np.array(dev_X)
```

```
In [91]: dev_X_test.shape
```

```
Out[91]: (1101,)
```

```
In [92]: dev_X_t = pad_sequences(dev_X_test, maxlen=267)
```

```
In [93]: dev_X_t.shape
```

```
Out[93]: (1101, 267)
```

```
In [ ]:
```

```
In [94]: dev_y = pd.get_dummies(dev['truth']).values
```

```
In [95]: dev_y
```

```
Out[95]: array([[0, 0, 0, 1, 0],  
                [0, 0, 1, 0, 0],  
                [0, 0, 0, 1, 0],  
                ...,  
                [0, 1, 0, 0, 0],  
                [0, 0, 1, 0, 0],  
                [0, 1, 0, 0, 0]], dtype=uint8)
```

```
In [ ]:
```

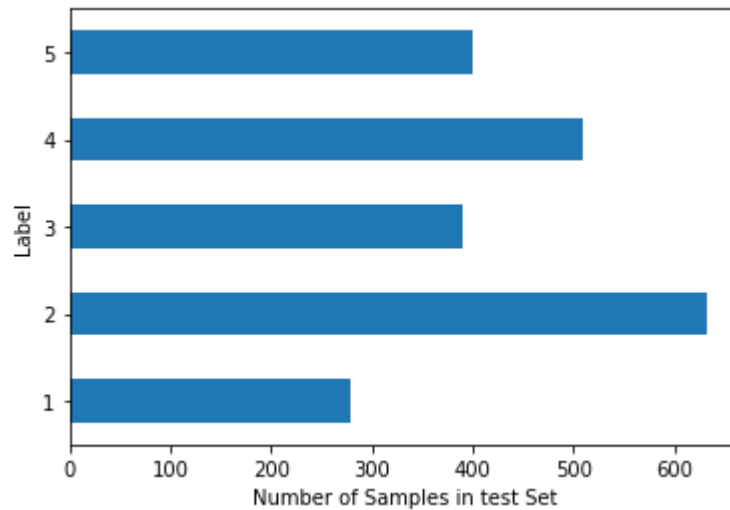
```
In [ ]:
```

```
In [ ]:
```

In [ ]:

## Test Data Visualizations

```
In [96]: ax = test['truth'].value_counts(sort=False).plot(kind='barh')
ax.set_xlabel("Number of Samples in test Set")
ax.set_ylabel("Label")
plt.show()
```



In [ ]:

```
In [97]: data_test = test['text']
```

```
In [98]: test_X = list()
```

```
In [99]: for i in range(len(test)):
integer_encoded = [ord(char) for char in data_test[i]]
test_X.append(integer_encoded)
```

```
In [100]: test_X_2 = np.array(test_X)
```

```
In [101]: test_X_2.shape
```

```
Out[101]: (2210,)
```

```
In [102]: test_X_t = pad_sequences(test_X_2, maxlen=267)
```

```
In [103]: test_X_t.shape
```

```
Out[103]: (2210, 267)
```

```
In [ ]:
```

```
In [104]: test_y = pd.get_dummies(test['truth']).values
```

```
In [105]: test_y
```

```
Out[105]: array([[0, 0, 1, 0, 0],
                 [0, 0, 0, 1, 0],
                 [0, 0, 0, 0, 1],
                 ...,
                 [0, 0, 0, 0, 1],
                 [0, 0, 0, 1, 0],
                 [1, 0, 0, 0, 0]], dtype=uint8)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

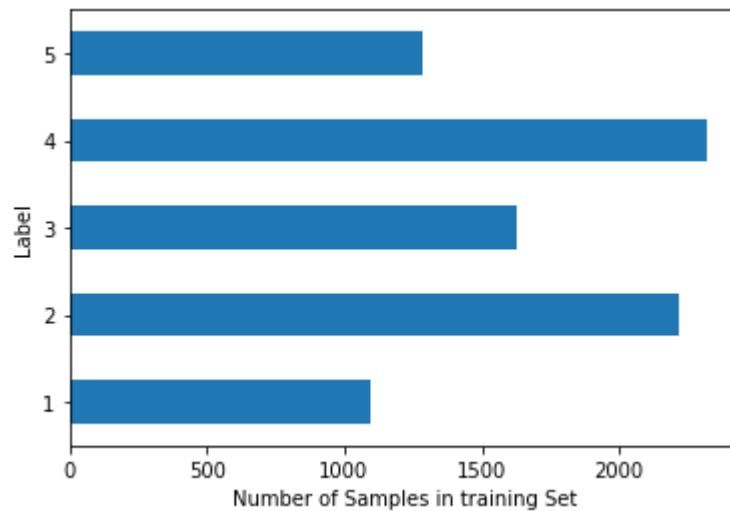


In [ ]:

## Training Data Visualizations

In [106]:

```
ax = df['truth'].value_counts(sort=False).plot(kind='barh')
ax.set_xlabel("Number of Samples in training Set")
ax.set_ylabel("Label")
plt.show()
```

In [277]: `y = df['truth'].values`In [278]: `y`

Out[278]: [4, 5, 4, 3, 4, ..., 1, 2, 4, 1, 2]  
Length: 8544  
Categories (5, int64): [1, 2, 3, 4, 5]

In [ ]:

```
In [109]: y_train = np.array(y)
```

```
In [276]: y_train
```

```
Out[276]: array([[0, 0, 0, 1, 0],  
                [0, 0, 1, 0, 0],  
                [0, 1, 0, 0, 0],  
                ...,  
                [0, 0, 1, 0, 0],  
                [0, 0, 0, 1, 0],  
                [0, 0, 1, 0, 0]], dtype=uint8)
```

```
In [111]: y_train[0]
```

```
Out[111]: 4
```

```
In [112]: y_train.shape
```

```
Out[112]: (8544,)
```

```
In [ ]:
```

```
In [113]: y_t = pd.get_dummies(df['truth']).values
```

```
In [114]: y_t
```

```
Out[114]: array([[0, 0, 0, 1, 0],  
                [0, 0, 0, 0, 1],  
                [0, 0, 0, 1, 0],  
                ...,  
                [0, 0, 0, 1, 0],  
                [1, 0, 0, 0, 0],  
                [0, 1, 0, 0, 0]], dtype=uint8)
```

```
In [115]: len(y_t)
```

```
Out[115]: 8544
```

In [ ]:

In [ ]:

```
In [116]: from numpy import argmax
# define input string
#data = 'hello world'
data = df['text']
print(data[0])
```

The Rock is destined to be the 21st Century 's new `` Conan '' and that he 's going to make a splash even greater than Arnold Schwarzenegger , Jean-Claud Van Damme or Steven Segal .

In [ ]:

In [ ]:

In [ ]:

In [ ]:

```
In [117]: X = list()
```

```
In [118]: #X.append(onehot_encoded)

#X.append(integer_encoded)
```

```
In [119]: len(data)
```

```
Out[119]: 8544
```

In [ ]:

In [ ]:

```
In [120]: for i in range(len(data)):
           integer_encoded = [ord(char) for char in data[i]]
           X.append(integer_encoded)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [121]: len(X[1])
```

```
Out[121]: 226
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]: X
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [123]: #X = np.array(onehot_encoded)
```

```
In [124]: X_train = np.array(X)
```

```
In [ ]:
```

```
In [ ]: X_train
```

```
In [126]: X_train.shape
```

```
Out[126]: (8544,)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [127]: # invert encoding  
#inverted = int_to_char[argmax(onehot_encoded[0])]   
#print(inverted)
```

```
In [ ]:
```

```
In [128]: X_train.shape
```

```
Out[128]: (8544,)
```

```
In [129]: X_t = pad_sequences(X_train)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [130]: X_train, X_test, y_train, y_test = train_test_split(X_t, y_t, test_size=0.2, random_state=0)
```

```
In [ ]:
```

```
In [ ]:
```

## Create vanilla RNN model

```
In [287]: model_rnn = tf.keras.Sequential()  
model_rnn.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))
```

```
In [288]: #model_rnn.add(layers.Dropout(0.3))
```

```
In [289]: model_rnn.add(layers.SimpleRNN(256))  
model_rnn.add(layers.Dense(5, activation='softmax'))
```

```
In [290]: model_rnn.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])  
model_rnn.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
=====		
embedding_6 (Embedding)	(None, 267, 256)	1280000
=====		
simple_rnn_7 (SimpleRNN)	(None, 256)	131328
=====		
dense_5 (Dense)	(None, 5)	1285
=====		
Total params: 1,412,613		
Trainable params: 1,412,613		
Non-trainable params: 0		
=====		

```
In [ ]:
```

```
In [291]: batch_size = 32  
epochs = 10
```

```
In [ ]:
```

```
In [292]: # Clear any logs from previous runs  
# !rm -rf ./logs/
```

```
In [ ]:
```

```
In [293]: log_dir_1="logs_1/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
In [294]: tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_1, histogram_freq=1)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [295]: #model_rnn.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_s
```

```
model_rnn.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_si
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/10

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (0.421986). Check your callbacks.

6835/6835 - 29s - loss: 1.6206 - accuracy: 0.2506 - val\_loss: 1.5783 - val\_accuracy: 0.2709

Epoch 2/10

6835/6835 - 27s - loss: 1.6193 - accuracy: 0.2600 - val\_loss: 1.5777 - val\_accuracy: 0.2709

Epoch 3/10

6835/6835 - 27s - loss: 1.6368 - accuracy: 0.2506 - val\_loss: 1.6156 - val\_accuracy: 0.2657

Epoch 4/10

6835/6835 - 26s - loss: 1.5861 - accuracy: 0.2587 - val\_loss: 1.5774 - val\_accuracy: 0.2697

Epoch 5/10

6835/6835 - 26s - loss: 1.5891 - accuracy: 0.2648 - val\_loss: 1.5862 - val\_accuracy: 0.2697

Epoch 6/10

6835/6835 - 26s - loss: 1.5865 - accuracy: 0.2713 - val\_loss: 1.5817 - val\_accuracy: 0.2703

Epoch 7/10

6835/6835 - 26s - loss: 1.5854 - accuracy: 0.2679 - val\_loss: 1.5842 - val\_accuracy: 0.2709

Epoch 8/10

6835/6835 - 26s - loss: 1.5814 - accuracy: 0.2614 - val\_loss: 1.5969 - val\_accuracy: 0.2715

Epoch 9/10

6835/6835 - 26s - loss: 1.5902 - accuracy: 0.2680 - val\_loss: 1.5852 - val\_accuracy: 0.2738

Epoch 10/10

6835/6835 - 26s - loss: 1.5839 - accuracy: 0.2625 - val\_loss: 1.5779 - val\_accuracy: 0.2709

```
Out[295]: <tensorflow.python.keras.callbacks.History at 0x7f15400fb410>
```

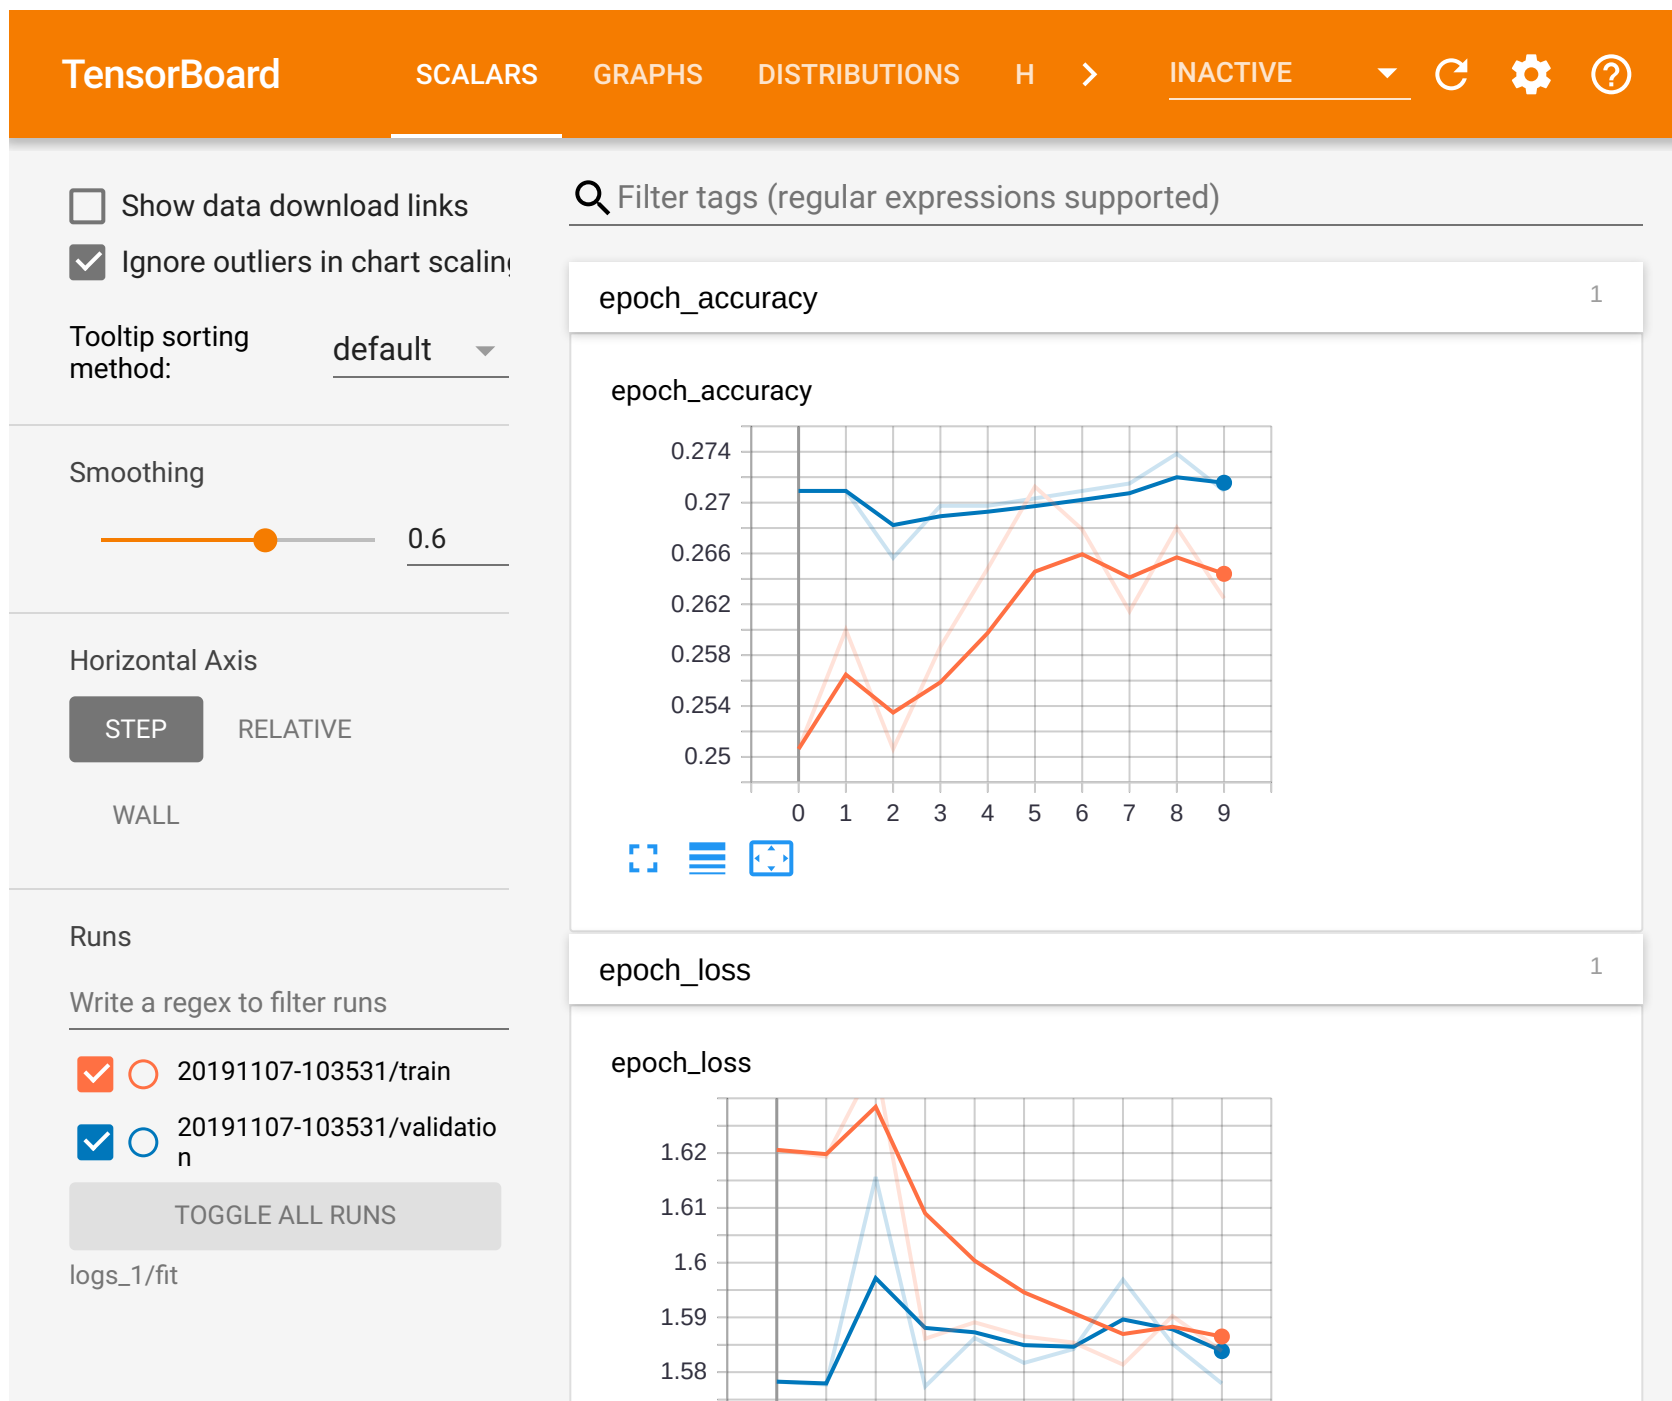
In [ ]:

## Tensorboard Visualizations

In [2]: `%load_ext tensorboard`



```
In [169]: %tensorboard --logdir logs_1/fit
```





In [ ]:

In [ ]:

In [ ]:

### Accuracy of vanilla rnn model on dev subset

In [ ]: `score, acc = model_rnn.evaluate(dev_X_t, dev_y)`In [343]: `acc`

Out[343]: 0.26430517

### Accuracy of vanilla rnn on dev subset = 26.43%

In [ ]:

In [ ]:

In [ ]:

In [ ]:

### Create vanilla RNN model - 2

```
In [397]: model_rnn_2 = tf.keras.Sequential()
model_rnn_2.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))

model_rnn_2.add(layers.SimpleRNN(256))
model_rnn_2.add(layers.Dense(5, activation='softmax'))

adam = tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, amsgrad=False)
#sgd = SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)

model_rnn_2.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
#model_rnn.summary()

batch_size = 32
epochs = 20

log_dir_5="logs_5/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_5, histogram_freq=1)

model_rnn_2.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/20

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (0.442880). Check your callbacks.

6835/6835 - 28s - loss: 1.6073 - accuracy: 0.2534 - val\_loss: 1.6004 - val\_accuracy: 0.2703

Epoch 2/20

6835/6835 - 26s - loss: 1.5903 - accuracy: 0.2660 - val\_loss: 1.5640 - val\_accuracy: 0.2703

Epoch 3/20

6835/6835 - 26s - loss: 1.6017 - accuracy: 0.2614 - val\_loss: 1.5946 - val\_accuracy: 0.1861

Epoch 4/20

6835/6835 - 26s - loss: 1.5865 - accuracy: 0.2631 - val\_loss: 1.6006 - val\_accuracy: 0.1861

Epoch 5/20

6835/6835 - 26s - loss: 1.5823 - accuracy: 0.2718 - val\_loss: 1.5997 - val\_accuracy: 0.2709

Epoch 6/20

6835/6835 - 26s - loss: 1.5830 - accuracy: 0.2691 - val\_loss: 1.5618 - val\_accuracy: 0.2738

Epoch 7/20

6835/6835 - 27s - loss: 1.5866 - accuracy: 0.2658 - val\_loss: 1.6199 - val\_accuracy: 0.1270

Epoch 8/20

6835/6835 - 26s - loss: 1.5873 - accuracy: 0.2604 - val\_loss: 1.6174 - val\_accuracy: 0.2738

```
Epoch 9/20
6835/6835 - 26s - loss: 1.5778 - accuracy: 0.2642 - val_loss: 1.5705 - val_accuracy: 0.2750
Epoch 10/20
6835/6835 - 26s - loss: 1.5817 - accuracy: 0.2729 - val_loss: 1.5705 - val_accuracy: 0.2697
Epoch 11/20
6835/6835 - 26s - loss: 1.5857 - accuracy: 0.2645 - val_loss: 1.5676 - val_accuracy: 0.2762
Epoch 12/20
6835/6835 - 26s - loss: 1.5901 - accuracy: 0.2606 - val_loss: 1.5588 - val_accuracy: 0.2721
Epoch 13/20
6835/6835 - 26s - loss: 1.5801 - accuracy: 0.2603 - val_loss: 1.5693 - val_accuracy: 0.2721
Epoch 14/20
6835/6835 - 26s - loss: 1.5867 - accuracy: 0.2576 - val_loss: 1.5688 - val_accuracy: 0.2733
Epoch 15/20
6835/6835 - 26s - loss: 1.5741 - accuracy: 0.2768 - val_loss: 1.5646 - val_accuracy: 0.2709
Epoch 16/20
6835/6835 - 26s - loss: 1.5921 - accuracy: 0.2711 - val_loss: 1.5761 - val_accuracy: 0.2721
Epoch 17/20
6835/6835 - 26s - loss: 1.5828 - accuracy: 0.2636 - val_loss: 1.5698 - val_accuracy: 0.2709
Epoch 18/20
6835/6835 - 26s - loss: 1.5804 - accuracy: 0.2629 - val_loss: 1.5748 - val_accuracy: 0.2703
Epoch 19/20
6835/6835 - 26s - loss: 1.5795 - accuracy: 0.2686 - val_loss: 1.5655 - val_accuracy: 0.2733
Epoch 20/20
6835/6835 - 26s - loss: 1.5825 - accuracy: 0.2753 - val_loss: 1.5611 - val_accuracy: 0.2715
```

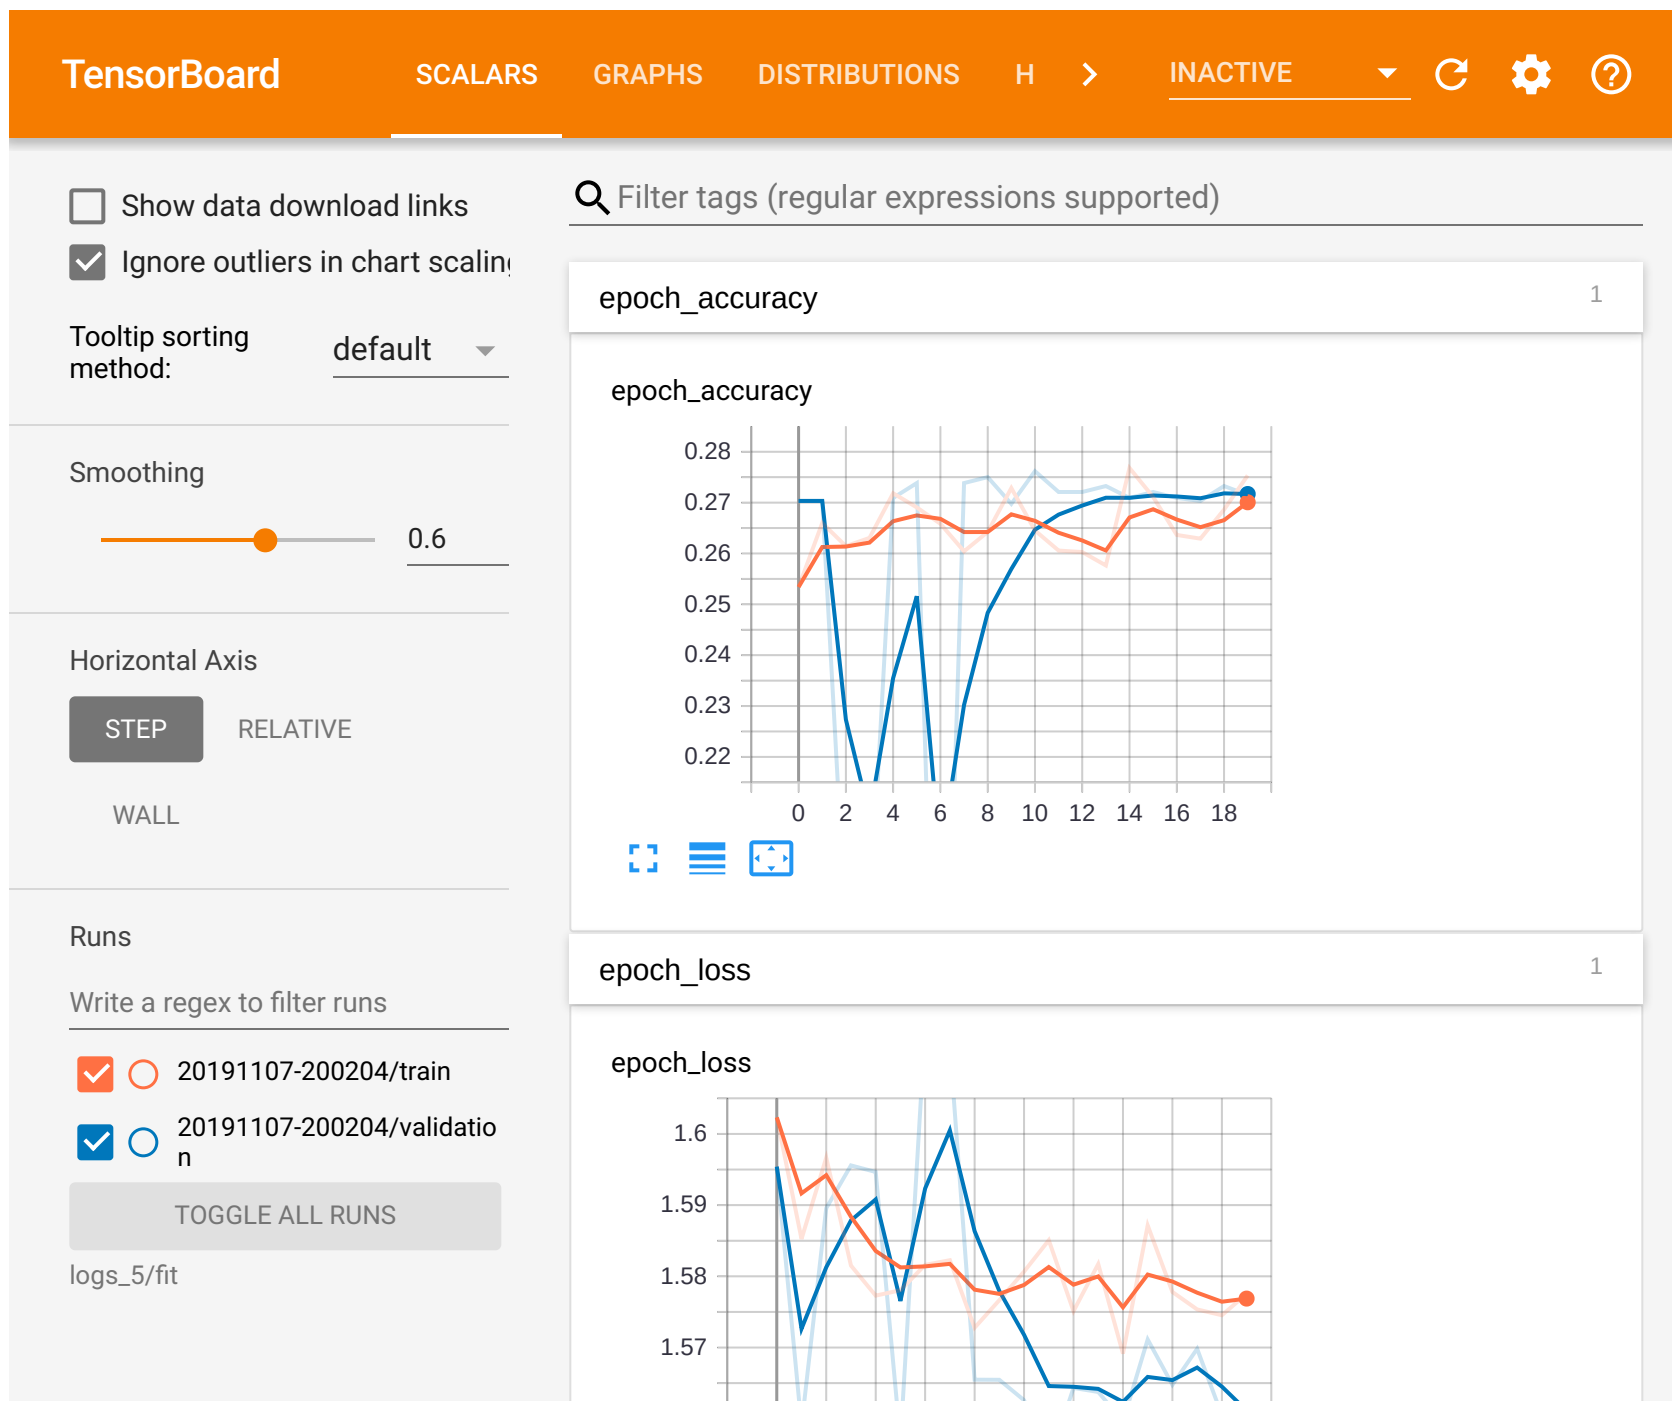
**Out[397]:** <tensorflow.python.keras.callbacks.History at 0x7f153a809210>


In [ ]:

In [ ]:

**With learning\_rate = 0.001**

```
In [170]: %tensorboard --logdir logs_5/fit
```





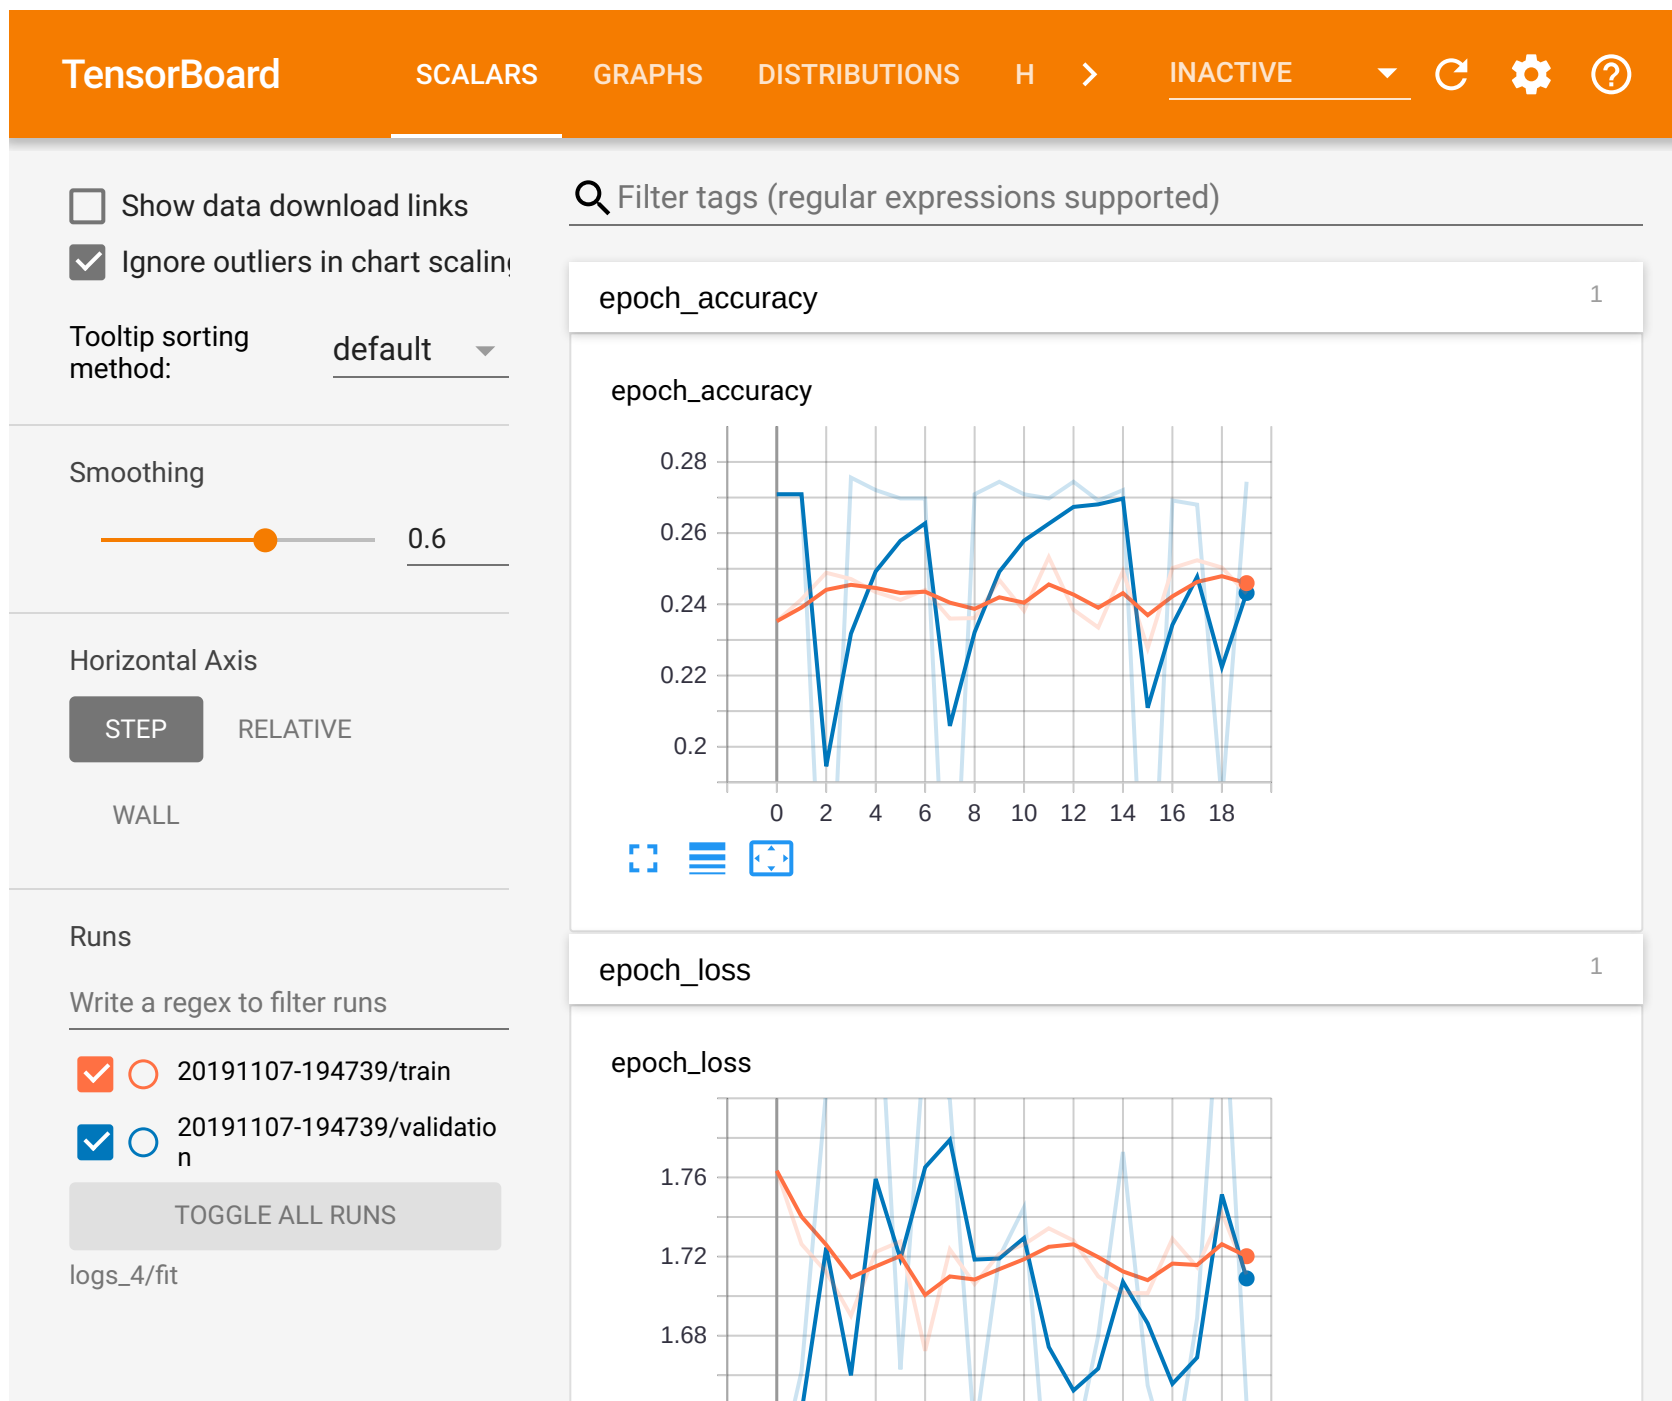
In [ ]:


In [ ]:

In [ ]:

**With learning\_rate = 0.01**

```
In [171]: %tensorboard --logdir logs_4/fit
```





In [ ]:

Create vanilla RNN model - 3



```
In [251]: model_rnn_3 = tf.keras.Sequential()
model_rnn_3.add(layers.Embedding(5000, 350, input_length=X_train.shape[1]))

model_rnn_3.add(layers.SimpleRNN(350))
model_rnn_3.add(layers.Dense(5, activation='softmax'))

adam = tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, amsgrad=False)
#sgd = SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)

model_rnn_3.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
#model_rnn.summary()

batch_size = 32
epochs = 20

log_dir_19="logs_19/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_19, histogram_freq=1)

model_rnn_3.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/20

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (0.440004). Check your callbacks.

6835/6835 - 44s - loss: 1.6198 - accuracy: 0.2493 - val\_loss: 1.5842 - val\_accuracy: 0.2703

Epoch 2/20

6835/6835 - 41s - loss: 1.6034 - accuracy: 0.2559 - val\_loss: 1.5649 - val\_accuracy: 0.2709

Epoch 3/20

6835/6835 - 41s - loss: 1.5930 - accuracy: 0.2613 - val\_loss: 1.6117 - val\_accuracy: 0.1861

Epoch 4/20

6835/6835 - 41s - loss: 1.5893 - accuracy: 0.2673 - val\_loss: 1.6076 - val\_accuracy: 0.1732

Epoch 5/20

6835/6835 - 41s - loss: 1.5954 - accuracy: 0.2582 - val\_loss: 1.5960 - val\_accuracy: 0.1861

Epoch 6/20

6835/6835 - 41s - loss: 1.5898 - accuracy: 0.2614 - val\_loss: 1.5620 - val\_accuracy: 0.2727

Epoch 7/20

6835/6835 - 41s - loss: 1.5871 - accuracy: 0.2623 - val\_loss: 1.5816 - val\_accuracy: 0.2703

Epoch 8/20

6835/6835 - 41s - loss: 1.5912 - accuracy: 0.2667 - val\_loss: 1.5778 - val\_accuracy: 0.2709

Epoch 9/20

```
6835/6835 - 41s - loss: 1.5873 - accuracy: 0.2604 - val_loss: 1.5837 - val_accuracy: 0.2697
Epoch 10/20
6835/6835 - 41s - loss: 1.5902 - accuracy: 0.2632 - val_loss: 1.5902 - val_accuracy: 0.2715
Epoch 11/20
6835/6835 - 41s - loss: 1.5956 - accuracy: 0.2632 - val_loss: 1.5835 - val_accuracy: 0.2709
Epoch 12/20
6835/6835 - 41s - loss: 1.5853 - accuracy: 0.2647 - val_loss: 1.5678 - val_accuracy: 0.2709
Epoch 13/20
6835/6835 - 41s - loss: 1.5883 - accuracy: 0.2679 - val_loss: 1.5810 - val_accuracy: 0.2791
Epoch 14/20
6835/6835 - 41s - loss: 1.5834 - accuracy: 0.2562 - val_loss: 1.5687 - val_accuracy: 0.2850
Epoch 15/20
6835/6835 - 41s - loss: 1.5777 - accuracy: 0.2736 - val_loss: 1.5724 - val_accuracy: 0.2709
Epoch 16/20
6835/6835 - 41s - loss: 1.5853 - accuracy: 0.2657 - val_loss: 1.5850 - val_accuracy: 0.2159
Epoch 17/20
6835/6835 - 41s - loss: 1.5907 - accuracy: 0.2666 - val_loss: 1.5771 - val_accuracy: 0.2709
Epoch 18/20
6835/6835 - 41s - loss: 1.5803 - accuracy: 0.2657 - val_loss: 1.5687 - val_accuracy: 0.2861
Epoch 19/20
6835/6835 - 41s - loss: 1.5808 - accuracy: 0.2676 - val_loss: 1.6080 - val_accuracy: 0.1843
Epoch 20/20
6835/6835 - 41s - loss: 1.5847 - accuracy: 0.2672 - val_loss: 1.5827 - val_accuracy: 0.2499
```

Out[251]: <tensorflow.python.keras.callbacks.History at 0x7f3d457adad0>

In [262]: %tensorboard --logdir logs\_19/fit

```
ERROR: Failed to launch TensorBoard (exited with 255).
Contents of stderr:
E1112 00:12:30.043277 140338601187136 program.py:226] TensorBoard could not bind to any port around
6006 (tried 10 times)
ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)
```

In [ ]:

## Create vanilla RNN model - 4 - Double Layer size

```
In [194]: model_rnn_4 = tf.keras.Sequential()
model_rnn_4.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))

model_rnn_4.add(layers.SimpleRNN(256, return_sequences=True))

model_rnn_4.add(layers.SimpleRNN(256))

model_rnn_4.add(layers.Dense(5, activation='softmax'))

adam = tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, amsgrad=False)
#sgd = SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)

model_rnn_4.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
#model_rnn.summary()

batch_size = 32
epochs = 20

log_dir_14="logs_14/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_14, histogram_freq=1)

model_rnn_4.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/20

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (0.918587). Check your callbacks.

6835/6835 - 53s - loss: 1.7610 - accuracy: 0.2136 - val\_loss: 1.5712 - val\_accuracy: 0.2657

Epoch 2/20

6835/6835 - 49s - loss: 1.6294 - accuracy: 0.2530 - val\_loss: 1.5805 - val\_accuracy: 0.2662

Epoch 3/20

6835/6835 - 49s - loss: 1.5864 - accuracy: 0.2591 - val\_loss: 1.5716 - val\_accuracy: 0.2662

Epoch 4/20

6835/6835 - 51s - loss: 1.5836 - accuracy: 0.2626 - val\_loss: 1.5884 - val\_accuracy: 0.2703

Epoch 5/20

6835/6835 - 50s - loss: 1.5904 - accuracy: 0.2634 - val\_loss: 1.5725 - val\_accuracy: 0.2662

Epoch 6/20

6835/6835 - 51s - loss: 1.5873 - accuracy: 0.2638 - val\_loss: 1.5823 - val\_accuracy: 0.2662

Epoch 7/20

```
6835/6835 - 49s - loss: 1.5872 - accuracy: 0.2617 - val_loss: 1.5919 - val_accuracy: 0.2703
Epoch 8/20
6835/6835 - 49s - loss: 1.5893 - accuracy: 0.2655 - val_loss: 1.5754 - val_accuracy: 0.2662
Epoch 9/20
6835/6835 - 49s - loss: 1.5906 - accuracy: 0.2616 - val_loss: 1.5897 - val_accuracy: 0.2662
Epoch 10/20
6835/6835 - 50s - loss: 1.5841 - accuracy: 0.2636 - val_loss: 1.5692 - val_accuracy: 0.2703
Epoch 11/20
6835/6835 - 49s - loss: 1.5837 - accuracy: 0.2597 - val_loss: 1.5797 - val_accuracy: 0.2662
Epoch 12/20
6835/6835 - 49s - loss: 1.5839 - accuracy: 0.2753 - val_loss: 1.5697 - val_accuracy: 0.2703
Epoch 13/20
6835/6835 - 49s - loss: 1.5925 - accuracy: 0.2676 - val_loss: 1.5914 - val_accuracy: 0.1861
Epoch 14/20
6835/6835 - 49s - loss: 1.5848 - accuracy: 0.2553 - val_loss: 1.5916 - val_accuracy: 0.2662
Epoch 15/20
6835/6835 - 49s - loss: 1.5901 - accuracy: 0.2688 - val_loss: 1.5811 - val_accuracy: 0.2703
Epoch 16/20
6835/6835 - 49s - loss: 1.5884 - accuracy: 0.2632 - val_loss: 1.6444 - val_accuracy: 0.2662
Epoch 17/20
6835/6835 - 50s - loss: 1.5871 - accuracy: 0.2641 - val_loss: 1.5696 - val_accuracy: 0.2703
Epoch 18/20
6835/6835 - 51s - loss: 1.5979 - accuracy: 0.2506 - val_loss: 1.5639 - val_accuracy: 0.2662
Epoch 19/20
6835/6835 - 51s - loss: 1.5939 - accuracy: 0.2528 - val_loss: 1.5707 - val_accuracy: 0.2703
Epoch 20/20
6835/6835 - 51s - loss: 1.5829 - accuracy: 0.2606 - val_loss: 1.5771 - val_accuracy: 0.2703
```

Out[194]: <tensorflow.python.keras.callbacks.History at 0x7f3db5fd2e50>

## Doubling the layer of vanilla RNN does not improve validation accuracy

In [ ]:

In [ ]:

```
In [263]: %tensorboard --logdir logs_14/fit
```

```
ERROR: Failed to launch TensorBoard (exited with 255).  
Contents of stderr:  
E1112 00:12:55.089560 140568162531136 program.py:226] TensorBoard could not bind to any port around  
6006 (tried 10 times)  
ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)
```

```
In [ ]:
```

```
In [ ]:
```

## Create LSTM Model

```
In [298]: # Clear any logs from previous runs  
          #!rm -rf ./logs/
```

```
In [219]: log_dir_18="logs_18/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
In [220]: tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_18, histogram_freq=1)
```

```
In [ ]:
```

```
In [221]: model_lstm = tf.keras.Sequential()  
          model_lstm.add(layers.Embedding(5000, 500, input_length=X_train.shape[1]))
```

```
In [222]: model_lstm.add(layers.Dropout(0.5))
```

```
In [223]: model_lstm.add(layers.LSTM(500, dropout=0.5, recurrent_dropout=0.5))  
          model_lstm.add(layers.Dense(5, activation='softmax'))
```

```
In [ ]:
```

```
In [224]: model_lstm.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])  
model_lstm.summary()
```

Model: "sequential\_24"

Layer (type)	Output Shape	Param #
=====		
embedding_7 (Embedding)	(None, 267, 500)	25000000
-----		
dropout_4 (Dropout)	(None, 267, 500)	0
-----		
lstm_28 (LSTM)	(None, 500)	2002000
-----		
dense_13 (Dense)	(None, 5)	2505
=====		
Total params: 4,504,505		
Trainable params: 4,504,505		
Non-trainable params: 0		
-----		

In [ ]:

In [ ]:

```
In [225]: batch_size = 32  
epochs = 15
```

```
In [226]: model_lstm.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_s
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/15

6835/6835 - 276s - loss: 1.5758 - accuracy: 0.2742 - val\_loss: 1.5610 - val\_accuracy: 0.2756

Epoch 2/15

6835/6835 - 272s - loss: 1.5708 - accuracy: 0.2683 - val\_loss: 1.5681 - val\_accuracy: 0.2756

Epoch 3/15

6835/6835 - 277s - loss: 1.5636 - accuracy: 0.2907 - val\_loss: 1.5733 - val\_accuracy: 0.2733

Epoch 4/15

6835/6835 - 272s - loss: 1.5649 - accuracy: 0.2803 - val\_loss: 1.5561 - val\_accuracy: 0.2779

Epoch 5/15

6835/6835 - 271s - loss: 1.5614 - accuracy: 0.2806 - val\_loss: 1.5636 - val\_accuracy: 0.2750

Epoch 6/15

6835/6835 - 270s - loss: 1.5574 - accuracy: 0.2907 - val\_loss: 1.5569 - val\_accuracy: 0.2902

Epoch 7/15

6835/6835 - 271s - loss: 1.5593 - accuracy: 0.2910 - val\_loss: 1.5576 - val\_accuracy: 0.2762

Epoch 8/15

6835/6835 - 271s - loss: 1.5548 - accuracy: 0.2926 - val\_loss: 1.5567 - val\_accuracy: 0.2844

Epoch 9/15

6835/6835 - 272s - loss: 1.5549 - accuracy: 0.2854 - val\_loss: 1.5529 - val\_accuracy: 0.2978

Epoch 10/15

6835/6835 - 272s - loss: 1.5553 - accuracy: 0.2970 - val\_loss: 1.5558 - val\_accuracy: 0.2978

Epoch 11/15

6835/6835 - 271s - loss: 1.5503 - accuracy: 0.2977 - val\_loss: 1.5567 - val\_accuracy: 0.2867

Epoch 12/15

6835/6835 - 271s - loss: 1.5450 - accuracy: 0.2945 - val\_loss: 1.5562 - val\_accuracy: 0.2774

Epoch 13/15

6835/6835 - 271s - loss: 1.5451 - accuracy: 0.3001 - val\_loss: 1.5622 - val\_accuracy: 0.2756

Epoch 14/15

6835/6835 - 271s - loss: 1.5409 - accuracy: 0.3050 - val\_loss: 1.5561 - val\_accuracy: 0.3037

Epoch 15/15

6835/6835 - 274s - loss: 1.5380 - accuracy: 0.3078 - val\_loss: 1.5553 - val\_accuracy: 0.2908

```
Out[226]: <tensorflow.python.keras.callbacks.History at 0x7f3dbda0a610>
```

In [264]: `%tensorboard --logdir logs_18/fit`

```
ERROR: Failed to launch TensorBoard (exited with 255).
Contents of stderr:
E1112 00:13:29.759110 140098707699520 program.py:226] TensorBoard could not bind to any port around
6006 (tried 10 times)
ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)
```

In [228]: `loss, acc = model_lstm.evaluate(test_X_t, test_y, verbose=2)`  
`print("Test set Accuracy: {:.2f}%".format(100*acc))`

```
2210/1 - 29s - loss: 1.5446 - accuracy: 0.2588
Test set Accuracy: 25.88%
```

In [ ]:

In [ ]:

In [ ]:

In [218]: `loss, acc = model_lstm.evaluate(test_X_t, test_y, verbose=2)`  
`print("Test set Accuracy: {:.2f}%".format(100*acc))`

```
2210/1 - 9s - loss: 1.5518 - accuracy: 0.2683
Test set Accuracy: 26.83%
```

In [ ]:

In [265]: `%tensorboard --logdir logs_17/fit`

```
ERROR: Failed to launch TensorBoard (exited with 255).
Contents of stderr:
E1112 00:13:54.057028 139840611579712 program.py:226] TensorBoard could not bind to any port around
6006 (tried 10 times)
ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)
```

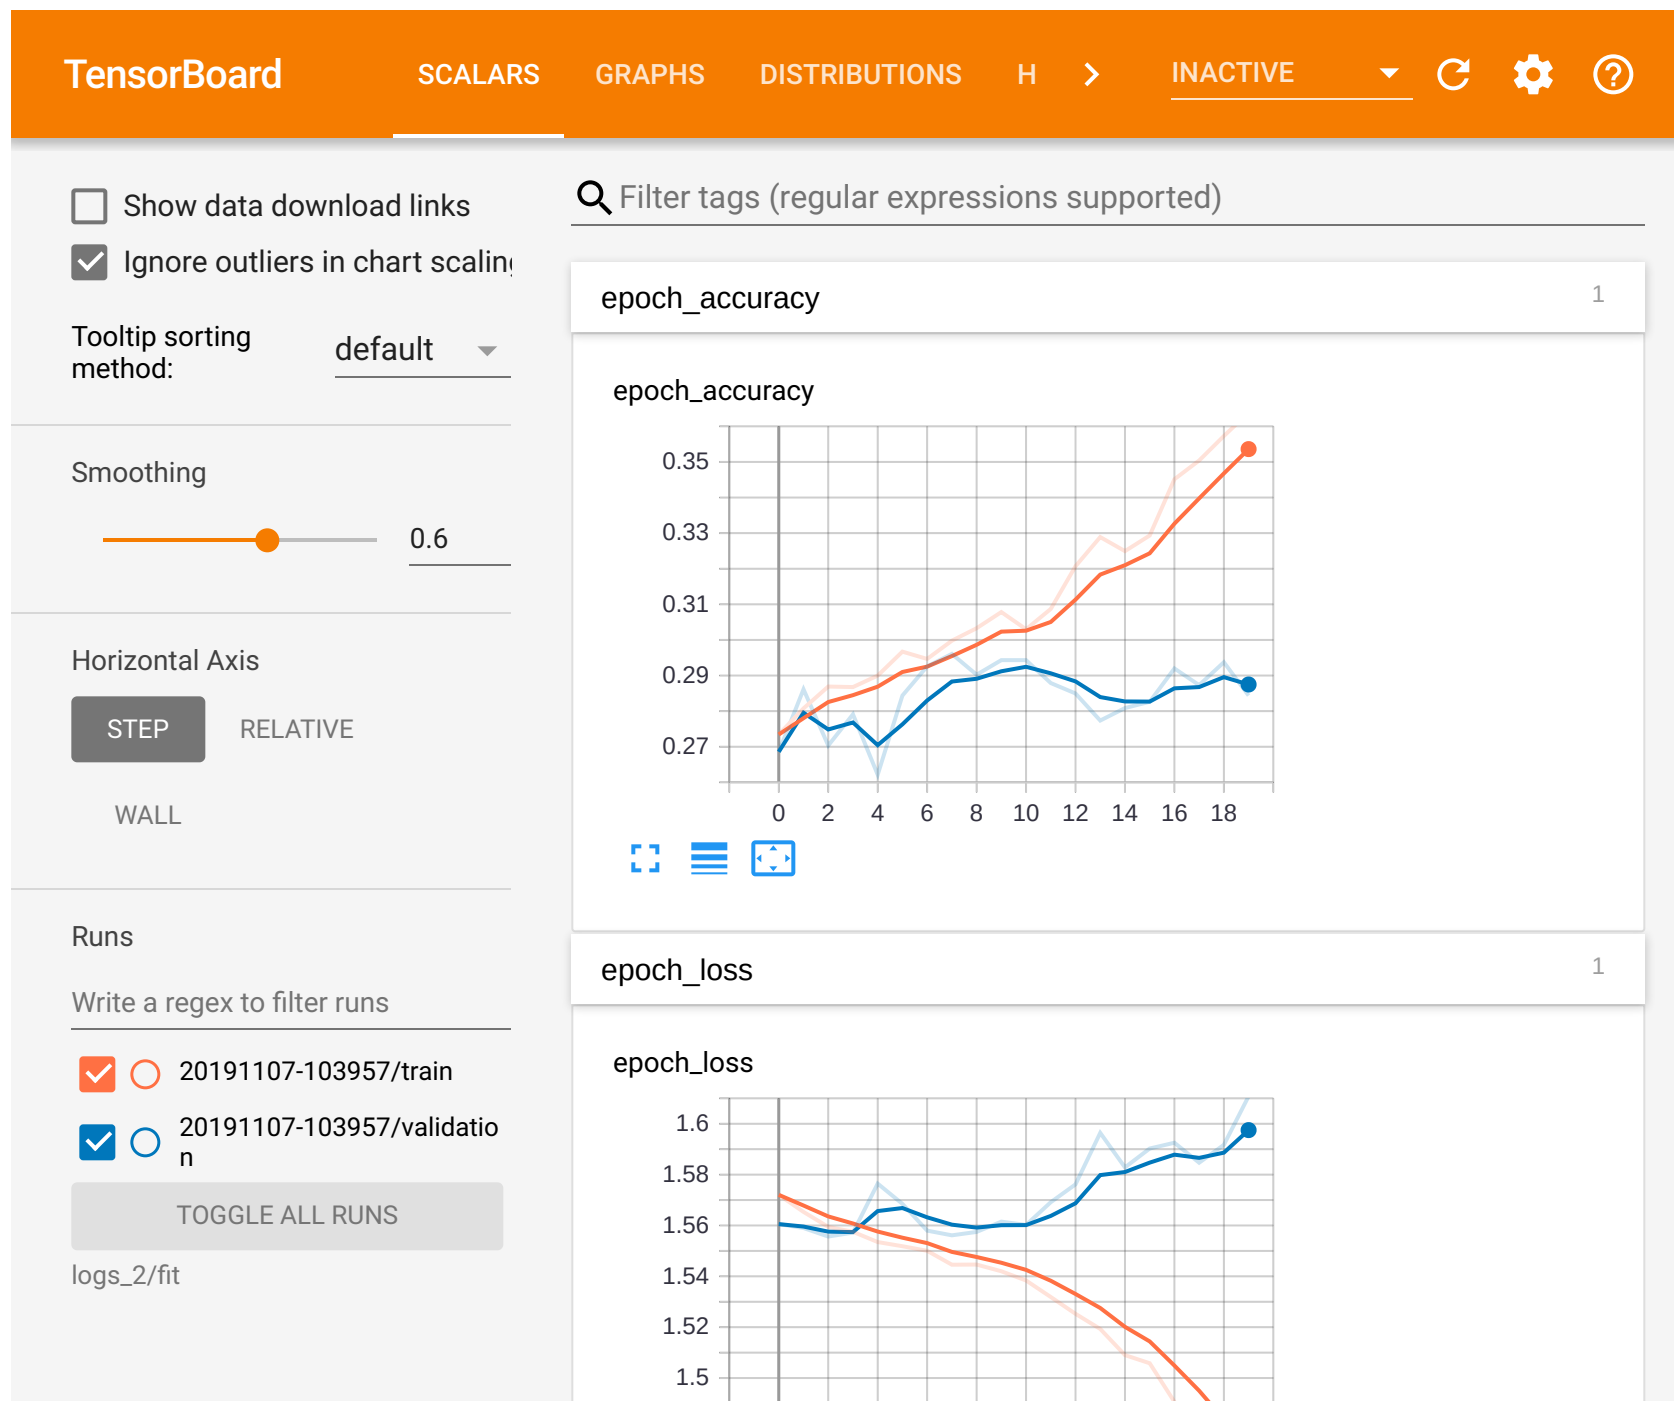
In [ ]:



In [ ]:

## LSTM Tensorboard Visualizations

```
In [173]: %tensorboard --logdir logs_2/fit
```





In [ ]:

### Accuracy of LSTM model on dev subset

In [ ]: `score, acc = model_lstm.evaluate(dev_X_t, dev_y)`

In [345]: `acc`

Out[345]: 0.2770209

### Accuracy of LSTM on dev subset = 27.70%

In [ ]:

In [ ]:

### Create LSTM Model - 2

```
In [401]: log_dir_7="logs_7/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
          tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_7, histogram_freq=1)
```

```
In [402]: model_lstm_2 = tf.keras.Sequential()
model_lstm_2.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))

model_lstm_2.add(layers.Dropout(0.3))

model_lstm_2.add(layers.LSTM(256, dropout=0.3, recurrent_dropout=0.2))
model_lstm_2.add(layers.Dense(5, activation='softmax'))

adam = tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, amsgrad=False)

model_lstm_2.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
#model_lstm.summary()

batch_size = 16
epochs = 100

model_lstm_2.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/100

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (0.927149). Check your callbacks.

6835/6835 - 124s - loss: 1.5718 - accuracy: 0.2713 - val\_loss: 1.5603 - val\_accuracy: 0.2803

Epoch 2/100

6835/6835 - 119s - loss: 1.5649 - accuracy: 0.2809 - val\_loss: 1.5618 - val\_accuracy: 0.2879

Epoch 3/100

6835/6835 - 119s - loss: 1.5611 - accuracy: 0.2799 - val\_loss: 1.5610 - val\_accuracy: 0.2756

Epoch 4/100

6835/6835 - 119s - loss: 1.5587 - accuracy: 0.2909 - val\_loss: 1.5599 - val\_accuracy: 0.2815

Epoch 5/100

6835/6835 - 119s - loss: 1.5556 - accuracy: 0.2938 - val\_loss: 1.5612 - val\_accuracy: 0.2768

Epoch 6/100

6835/6835 - 120s - loss: 1.5554 - accuracy: 0.2857 - val\_loss: 1.5584 - val\_accuracy: 0.2756

Epoch 7/100

6835/6835 - 119s - loss: 1.5474 - accuracy: 0.2951 - val\_loss: 1.5655 - val\_accuracy: 0.2709

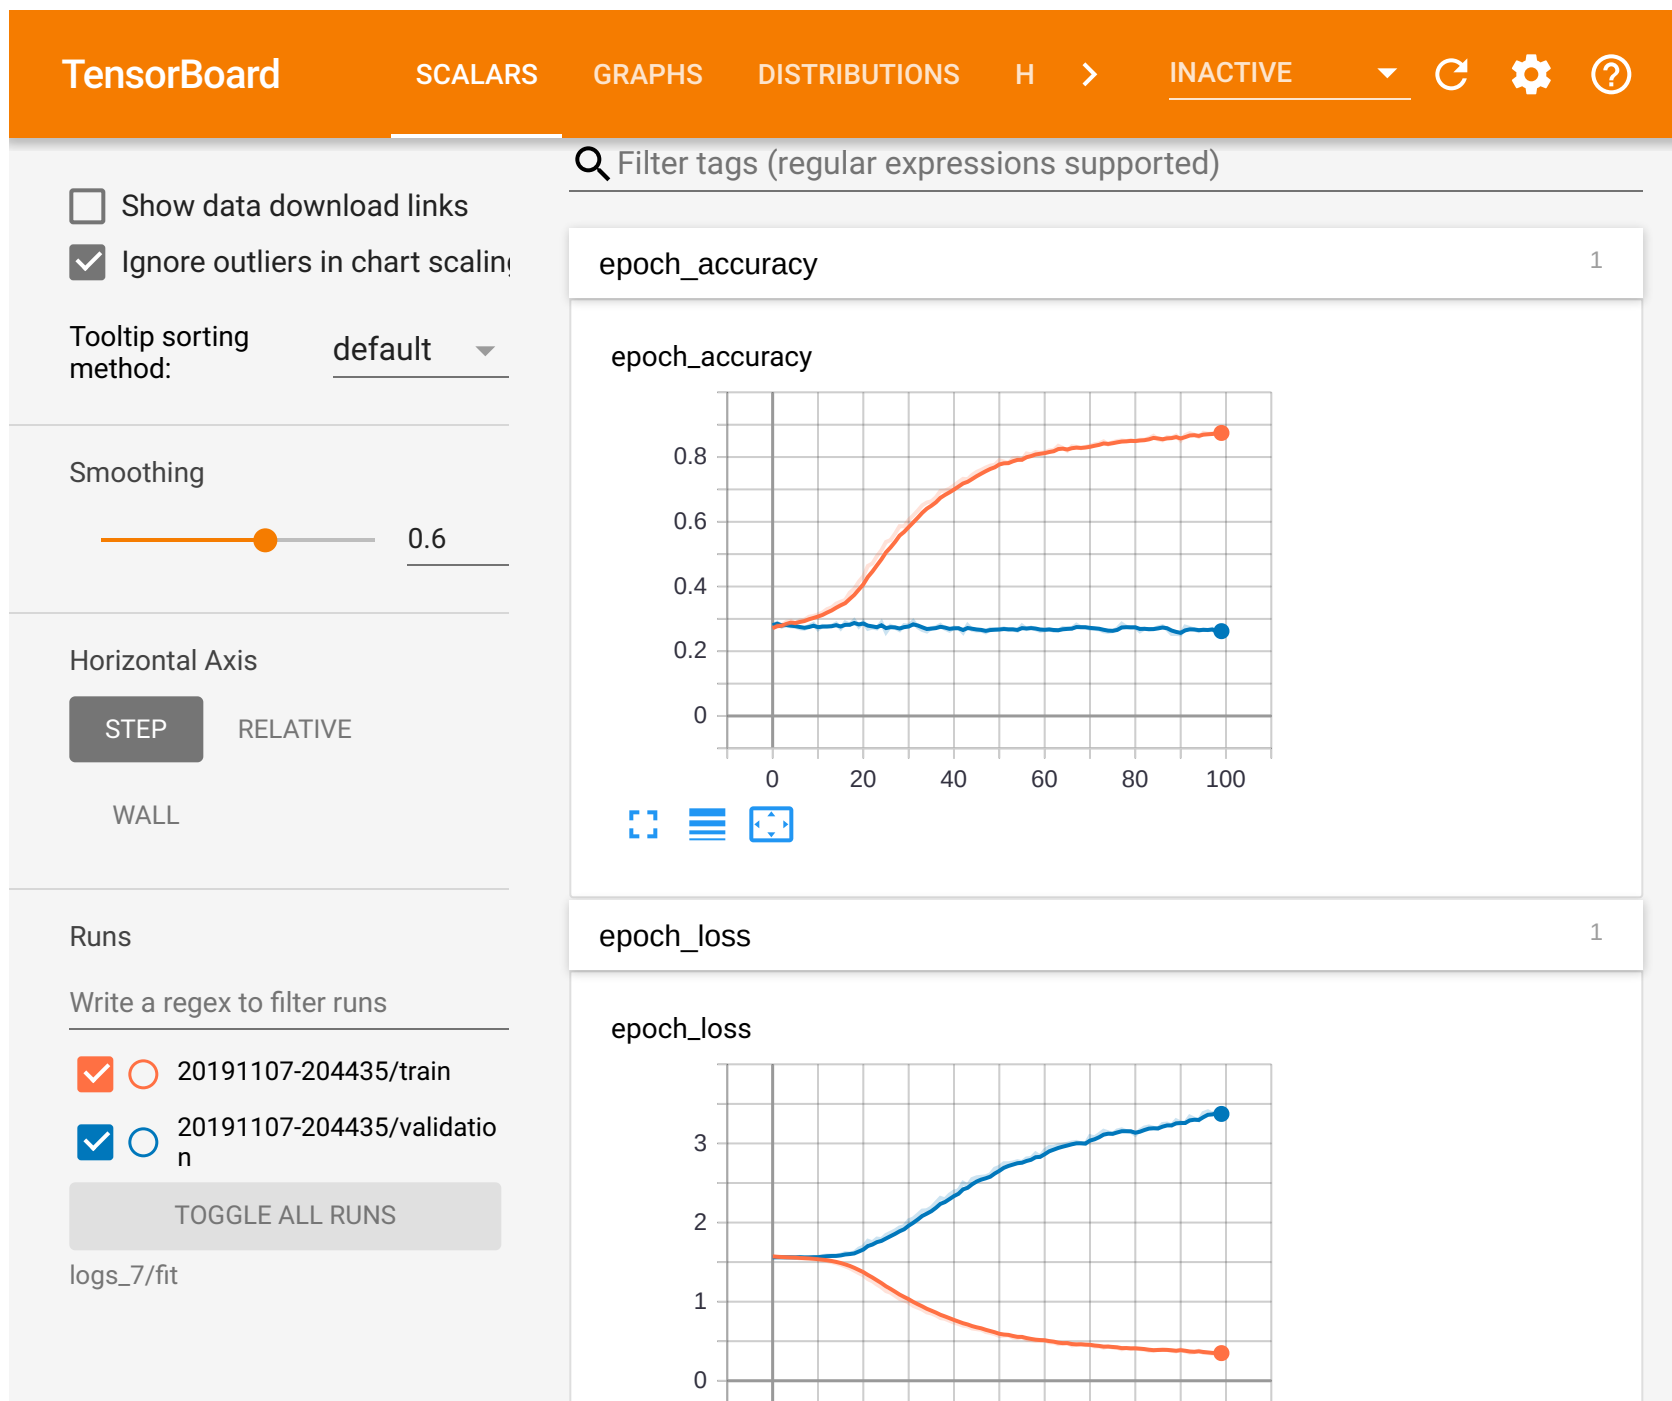
Epoch 8/100

6835/6835 - 119s - loss: 1.5454 - accuracy: 0.2973 - val\_loss: 1.5561 - val\_accuracy: 0.2692

Epoch 9/100

In [ ]:

```
In [174]: %tensorboard --logdir logs_7/fit
```





In [ ]:

**Doubling LSTM**

In [ ]:

```
In [196]: model_lstm_3 = tf.keras.Sequential()
model_lstm_3.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))

model_lstm_3.add(layers.Dropout(0.3))

model_lstm_3.add(layers.LSTM(256, return_sequences=True, dropout=0.3, recurrent_dropout=0.2))

model_lstm_3.add(layers.LSTM(256, dropout=0.3, recurrent_dropout=0.2))

model_lstm_3.add(layers.Dense(5, activation='softmax'))

adam = tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, amsgrad=False)

model_lstm_3.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])
#model_lstm.summary()

batch_size = 16
epochs = 10

log_dir_15="logs_15/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")

tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_15, histogram_freq=1)

model_lstm_3.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/10

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (1.967493). Check your callbacks.

6835/6835 - 247s - loss: 1.5763 - accuracy: 0.2598 - val\_loss: 1.5657 - val\_accuracy: 0.2668

Epoch 2/10

6835/6835 - 237s - loss: 1.5696 - accuracy: 0.2723 - val\_loss: 1.5665 - val\_accuracy: 0.2873

Epoch 3/10

6835/6835 - 237s - loss: 1.5683 - accuracy: 0.2783 - val\_loss: 1.5698 - val\_accuracy: 0.2715

Epoch 4/10

6835/6835 - 235s - loss: 1.5728 - accuracy: 0.2695 - val\_loss: 1.5652 - val\_accuracy: 0.2662



```
Epoch 5/10
6835/6835 - 235s - loss: 1.5691 - accuracy: 0.2699 - val_loss: 1.5630 - val_accuracy: 0.2697
Epoch 6/10
6835/6835 - 239s - loss: 1.5677 - accuracy: 0.2742 - val_loss: 1.5570 - val_accuracy: 0.2744
Epoch 7/10
6835/6835 - 236s - loss: 1.5667 - accuracy: 0.2740 - val_loss: 1.5583 - val_accuracy: 0.2727
Epoch 8/10
6835/6835 - 235s - loss: 1.5659 - accuracy: 0.2794 - val_loss: 1.5599 - val_accuracy: 0.2727
Epoch 9/10
6835/6835 - 235s - loss: 1.5654 - accuracy: 0.2780 - val_loss: 1.5584 - val_accuracy: 0.2733
Epoch 10/10
6835/6835 - 235s - loss: 1.5660 - accuracy: 0.2727 - val_loss: 1.5593 - val_accuracy: 0.2727
```

Out[196]: <tensorflow.python.keras.callbacks.History at 0x7f3f2c9c76d0>

In [ ]:

## Doubling LSTM layer does not improve validation accuracy

In [ ]:

In [268]: `%tensorboard --logdir logs_15/fit`

```
ERROR: Failed to launch TensorBoard (exited with 255).
Contents of stderr:
E1112 00:15:39.989207 140272832292672 program.py:226] TensorBoard could not bind to any port around
6006 (tried 10 times)
ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)
```

In [ ]:

In [ ]:

In [ ]:

## Create GRU model

In [ ]:

In [ ]:

```
In [238]: model_gru = tf.keras.Sequential()  
model_gru.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))
```

```
In [239]: model_gru.add(layers.GRU(256))  
model_gru.add(layers.Dense(5, activation='softmax'))
```

```
In [240]: model_gru.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])  
model_gru.summary()
```

Model: "sequential\_25"

Layer (type)	Output Shape	Param #
=====		
embedding_8 (Embedding)	(None, 267, 256)	1280000
-----		
gru_2 (GRU)	(None, 256)	394752
-----		
dense_14 (Dense)	(None, 5)	1285
=====		
Total params: 1,676,037		
Trainable params: 1,676,037		
Non-trainable params: 0		

```
In [241]: batch_size = 32  
epochs = 20
```

In [ ]:

```
In [242]: # Clear any logs from previous runs  
#!/rm -rf ./logs/
```

```
In [243]: log_dir_3="logs_3/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
In [244]: tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_3, histogram_freq=1)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [245]: model_gru.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_si
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/20

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (0.882006). Check your callbacks.

6835/6835 - 72s - loss: 1.5719 - accuracy: 0.2692 - val\_loss: 1.5626 - val\_accuracy: 0.2779

Epoch 2/20

6835/6835 - 68s - loss: 1.5653 - accuracy: 0.2742 - val\_loss: 1.5680 - val\_accuracy: 0.2727

Epoch 3/20

6835/6835 - 69s - loss: 1.5596 - accuracy: 0.2830 - val\_loss: 1.5621 - val\_accuracy: 0.2686

Epoch 4/20

6835/6835 - 69s - loss: 1.5551 - accuracy: 0.2939 - val\_loss: 1.5637 - val\_accuracy: 0.2744

Epoch 5/20

6835/6835 - 68s - loss: 1.5519 - accuracy: 0.2951 - val\_loss: 1.5624 - val\_accuracy: 0.2774

Epoch 6/20

6835/6835 - 68s - loss: 1.5481 - accuracy: 0.2985 - val\_loss: 1.5650 - val\_accuracy: 0.2762

Epoch 7/20

6835/6835 - 68s - loss: 1.5442 - accuracy: 0.3023 - val\_loss: 1.5559 - val\_accuracy: 0.2873

Epoch 8/20

6835/6835 - 68s - loss: 1.5398 - accuracy: 0.2999 - val\_loss: 1.5626 - val\_accuracy: 0.2978

Epoch 9/20

6835/6835 - 70s - loss: 1.5340 - accuracy: 0.3138 - val\_loss: 1.5642 - val\_accuracy: 0.2896

Epoch 10/20

6835/6835 - 68s - loss: 1.5210 - accuracy: 0.3108 - val\_loss: 1.5613 - val\_accuracy: 0.2972

Epoch 11/20

6835/6835 - 68s - loss: 1.5047 - accuracy: 0.3230 - val\_loss: 1.5706 - val\_accuracy: 0.2733

Epoch 12/20

6835/6835 - 68s - loss: 1.4695 - accuracy: 0.3501 - val\_loss: 1.5598 - val\_accuracy: 0.2949

Epoch 13/20

6835/6835 - 68s - loss: 1.4112 - accuracy: 0.3914 - val\_loss: 1.5836 - val\_accuracy: 0.2955

Epoch 14/20

6835/6835 - 68s - loss: 1.3208 - accuracy: 0.4388 - val\_loss: 1.6476 - val\_accuracy: 0.2815

Epoch 15/20

6835/6835 - 68s - loss: 1.1865 - accuracy: 0.4976 - val\_loss: 1.7373 - val\_accuracy: 0.2785

Epoch 16/20

6835/6835 - 68s - loss: 1.0149 - accuracy: 0.5890 - val\_loss: 1.8808 - val\_accuracy: 0.2955

Epoch 17/20

6835/6835 - 68s - loss: 0.8169 - accuracy: 0.6825 - val\_loss: 2.0533 - val\_accuracy: 0.3084

Epoch 18/20

6835/6835 - 68s - loss: 0.6009 - accuracy: 0.7760 - val\_loss: 2.4067 - val\_accuracy: 0.2850

Epoch 19/20

```
6835/6835 - 68s - loss: 0.4321 - accuracy: 0.8505 - val_loss: 2.6931 - val_accuracy: 0.2943
Epoch 20/20
6835/6835 - 69s - loss: 0.3065 - accuracy: 0.9002 - val_loss: 3.0015 - val_accuracy: 0.2932
```

Out[245]: <tensorflow.python.keras.callbacks.History at 0x7f3d54e10250>

In [ ]:

## GRU Tensorboard Visualizations

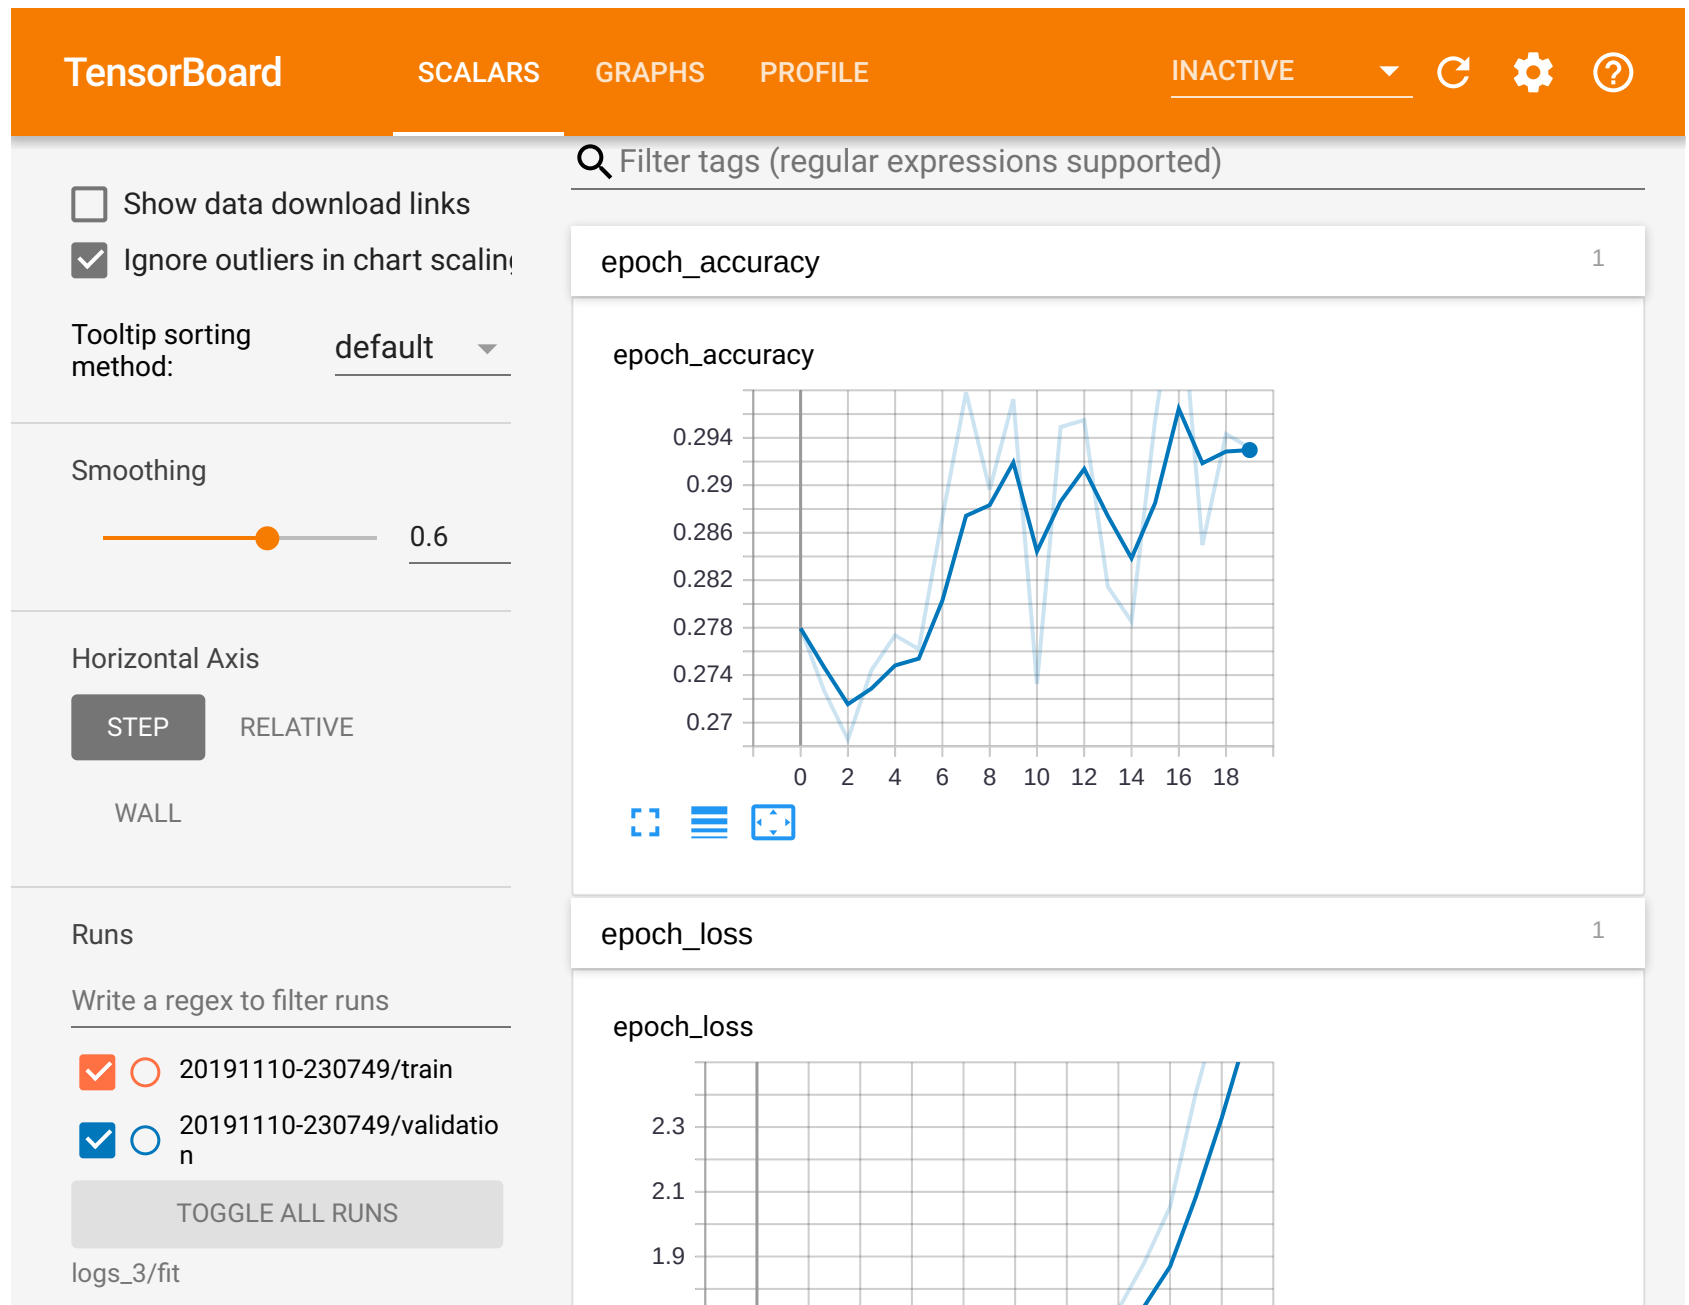
In [ ]:


In [269]: %load\_ext tensorboard

The tensorboard extension is already loaded. To reload it, use:  
%reload\_ext tensorboard

```
In [270]: %tensorboard --logdir logs_3/fit
```

Reusing TensorBoard on port 6015 (pid 20542), started 2 days, 15:33:50 ago. (Use '!kill 20542' to kill it.)





In [ ]:

In [ ]:

### Accuracy of GRU model on dev subset

In [ ]: `score, acc = model_gru.evaluate(dev_X_t, dev_y)`

In [347]: `acc`

Out[347]: 0.25613078

### Accuracy of GRU model on dev subset = 25.61%

In [ ]:

### Create GRU model - 2

```

In [418]: model_gru_2 = tf.keras.Sequential()
model_gru_2.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))

model_gru_2.add(layers.Dropout(0.3))

model_gru_2.add(layers.GRU(256))

model_gru_2.add(layers.Dropout(0.3))

model_gru_2.add(layers.Dense(5, activation='softmax'))

model_gru_2.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model_gru_2.summary()

batch_size = 32
epochs = 10

log_dir_9="logs_9/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_9, histogram_freq=1)

model_gru_2.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_

```

Model: "sequential\_27"

Layer (type)	Output Shape	Param #
embedding_26 (Embedding)	(None, 267, 256)	1280000
dropout_5 (Dropout)	(None, 267, 256)	0
gru_4 (GRU)	(None, 256)	394752
dropout_6 (Dropout)	(None, 256)	0



dense_25 (Dense)	(None, 5)	1285
------------------	-----------	------

---

Total params: 1,676,037  
 Trainable params: 1,676,037  
 Non-trainable params: 0

---

Train on 6835 samples, validate on 1709 samples

Epoch 1/10

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (0.891410). Check your callbacks.

6835/6835 - 74s - loss: 1.5734 - accuracy: 0.2658 - val\_loss: 1.5675 - val\_accuracy: 0.2744

Epoch 2/10

6835/6835 - 70s - loss: 1.5670 - accuracy: 0.2759 - val\_loss: 1.5578 - val\_accuracy: 0.2703

Epoch 3/10

6835/6835 - 70s - loss: 1.5620 - accuracy: 0.2811 - val\_loss: 1.5564 - val\_accuracy: 0.2896

Epoch 4/10

6835/6835 - 69s - loss: 1.5583 - accuracy: 0.2878 - val\_loss: 1.5621 - val\_accuracy: 0.2826

Epoch 5/10

6835/6835 - 70s - loss: 1.5564 - accuracy: 0.2876 - val\_loss: 1.5599 - val\_accuracy: 0.2844

Epoch 6/10

6835/6835 - 70s - loss: 1.5531 - accuracy: 0.2887 - val\_loss: 1.5581 - val\_accuracy: 0.2861

Epoch 7/10

6835/6835 - 70s - loss: 1.5491 - accuracy: 0.3031 - val\_loss: 1.5573 - val\_accuracy: 0.2756

Epoch 8/10

6835/6835 - 70s - loss: 1.5436 - accuracy: 0.3045 - val\_loss: 1.5545 - val\_accuracy: 0.2891

Epoch 9/10

6835/6835 - 70s - loss: 1.5370 - accuracy: 0.3058 - val\_loss: 1.5649 - val\_accuracy: 0.2779

Epoch 10/10

6835/6835 - 70s - loss: 1.5311 - accuracy: 0.3154 - val\_loss: 1.5549 - val\_accuracy: 0.2932

Out[418]: <tensorflow.python.keras.callbacks.History at 0x7f1389b57410>

In [271]: %tensorboard --logdir logs\_9/fit

ERROR: Failed to launch TensorBoard (exited with 255).

Contents of stderr:

E1112 00:16:19.276887 140082655811392 program.py:226] TensorBoard could not bind to any port around 6006 (tried 10 times)

ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)

In [ ]:

In [ ]:

### **Doubling GRU layer model**

In [ ]:

```
In [198]: model_gru_3 = tf.keras.Sequential()
model_gru_3.add(layers.Embedding(5000, 256, input_length=X_train.shape[1]))

model_gru_3.add(layers.Dropout(0.3))

model_gru_3.add(layers.GRU(256, return_sequences=True))

model_gru_3.add(layers.GRU(256))

model_gru_3.add(layers.Dropout(0.3))

model_gru_3.add(layers.Dense(5, activation='softmax'))

model_gru_3.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
#model_gru_3.summary()

batch_size = 32
epochs = 10

log_dir_16="logs_16/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_16, histogram_freq=1)

#model_gru_3.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)
model_gru_3.fit(X_t, y_t, validation_data=(dev_X_t, dev_y), epochs=epochs, batch_size=batch_size, ve
```

Train on 8544 samples, validate on 1101 samples

Epoch 1/10

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (1.201706). Check your callbacks.

8544/8544 - 160s - loss: 1.5742 - accuracy: 0.2697 - val\_loss: 1.5765 - val\_accuracy: 0.2543

Epoch 2/10

8544/8544 - 152s - loss: 1.5666 - accuracy: 0.2796 - val\_loss: 1.5808 - val\_accuracy: 0.2498

```
Epoch 3/10
8544/8544 - 152s - loss: 1.5653 - accuracy: 0.2773 - val_loss: 1.5706 - val_accuracy: 0.2625
Epoch 4/10
8544/8544 - 152s - loss: 1.5623 - accuracy: 0.2811 - val_loss: 1.5721 - val_accuracy: 0.2725
Epoch 5/10
8544/8544 - 152s - loss: 1.5599 - accuracy: 0.2786 - val_loss: 1.5771 - val_accuracy: 0.2480
Epoch 6/10
8544/8544 - 152s - loss: 1.5571 - accuracy: 0.2933 - val_loss: 1.5853 - val_accuracy: 0.2543
Epoch 7/10
8544/8544 - 152s - loss: 1.5530 - accuracy: 0.2921 - val_loss: 1.5823 - val_accuracy: 0.2707
Epoch 8/10
8544/8544 - 152s - loss: 1.5515 - accuracy: 0.2949 - val_loss: 1.5787 - val_accuracy: 0.2834
Epoch 9/10
8544/8544 - 152s - loss: 1.5456 - accuracy: 0.2990 - val_loss: 1.5703 - val_accuracy: 0.2570
Epoch 10/10
8544/8544 - 153s - loss: 1.5422 - accuracy: 0.2989 - val_loss: 1.5719 - val_accuracy: 0.2725
```

```
Out[198]: <tensorflow.python.keras.callbacks.History at 0x7f3dbd302090>
```

## Doubling GRU model layer does not improve validation accuracy

```
In [ ]:
```

```
In [199]: %tensorboard --logdir logs_16/fit
```

```
ERROR: Failed to launch TensorBoard (exited with 255).
Contents of stderr:
E1110 00:51:09.285538 139798023501632 program.py:226] TensorBoard could not bind to any port around
6006 (tried 10 times)
ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)
```

```
In [ ]:
```

```
In [ ]:
```

## Overfitting - with dropout=0.3

In [177]: `%tensorboard --logdir logs_8/fit`

ERROR: Failed to launch TensorBoard (exited with 255).

Contents of stderr:

E1109 08:42:37.262158 140054096025408 program.py:226] TensorBoard could not bind to any port around 6006 (tried 10 times)

ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)

In [ ]:

In [ ]:

In [134]: `X_train.shape`

Out[134]: (6835, 267)

In [ ]:

## **BiDIRECTIONAL RNN**

In [ ]:

```
In [158]: model_bid = tf.keras.Sequential()

model_bid.add(layers.Embedding(5000, 10, input_length=X_train.shape[1]))

model_bid.add(layers.Bidirectional(layers.LSTM(10, return_sequences=True)))

#model_bid.add(layers.Bidirectional(layers.LSTM(10, return_sequences=True), input_shape=(6835, 267)))

model_bid.add(layers.Bidirectional(layers.LSTM(10)))
model_bid.add(layers.Dense(5))
model_bid.add(layers.Activation('softmax'))
model_bid.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])

model_bid.summary()
```

Model: "sequential\_19"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 267, 10)	50000
bidirectional_17 (Bidirectional)	(None, 267, 20)	1680
bidirectional_18 (Bidirectional)	(None, 20)	2480
dense_8 (Dense)	(None, 5)	105
activation_7 (Activation)	(None, 5)	0
Total params: 54,265		
Trainable params: 54,265		
Non-trainable params: 0		

In [ ]:

```
In [166]: batch_size = 32
          epochs = 10

log_dir_13="logs_13/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir_13, histogram_freq=1)

# define checkpoint
checkpoint_path = "training_2/cp.ckpt"
checkpoint_dir = os.path.dirname(checkpoint_path)

# Create a callback that saves the model's weights
cp_callback = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path, monitor='accuracy', save_

model_bid.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_si
```

Train on 6835 samples, validate on 1709 samples

Epoch 1/10

WARNING:tensorflow:Method (on\_train\_batch\_end) is slow compared to the batch update (3.036925). Check your callbacks.

Epoch 00001: accuracy improved from -inf to 0.32963, saving model to training\_2/cp.ckpt

6835/6835 - 46s - loss: 1.5152 - accuracy: 0.3296 - val\_loss: 1.5677 - val\_accuracy: 0.2867

Epoch 2/10

Epoch 00002: accuracy did not improve from 0.32963

6835/6835 - 38s - loss: 1.5167 - accuracy: 0.3251 - val\_loss: 1.5604 - val\_accuracy: 0.2937

Epoch 3/10

Epoch 00003: accuracy did not improve from 0.32963

6835/6835 - 38s - loss: 1.5152 - accuracy: 0.3293 - val\_loss: 1.5642 - val\_accuracy: 0.2932

Epoch 4/10

Epoch 00004: accuracy improved from 0.32963 to 0.33021, saving model to training\_2/cp.ckpt

6835/6835 - 38s - loss: 1.5127 - accuracy: 0.3302 - val\_loss: 1.5741 - val\_accuracy: 0.2733

Epoch 5/10

Epoch 00005: accuracy improved from 0.33021 to 0.33270, saving model to training\_2/cp.ckpt

6835/6835 - 38s - loss: 1.5113 - accuracy: 0.3327 - val\_loss: 1.5642 - val\_accuracy: 0.2803

Epoch 6/10

Epoch 00006: accuracy improved from 0.33270 to 0.33899, saving model to training\_2/cp.ckpt  
6835/6835 - 38s - loss: 1.5082 - accuracy: 0.3390 - val\_loss: 1.5851 - val\_accuracy: 0.2744  
Epoch 7/10

Epoch 00007: accuracy did not improve from 0.33899  
6835/6835 - 38s - loss: 1.5074 - accuracy: 0.3386 - val\_loss: 1.5648 - val\_accuracy: 0.2844  
Epoch 8/10

Epoch 00008: accuracy did not improve from 0.33899  
6835/6835 - 38s - loss: 1.5098 - accuracy: 0.3302 - val\_loss: 1.5804 - val\_accuracy: 0.2885  
Epoch 9/10

Epoch 00009: accuracy did not improve from 0.33899  
6835/6835 - 38s - loss: 1.5092 - accuracy: 0.3333 - val\_loss: 1.5714 - val\_accuracy: 0.2744  
Epoch 10/10

Epoch 00010: accuracy did not improve from 0.33899  
6835/6835 - 38s - loss: 1.5120 - accuracy: 0.3260 - val\_loss: 1.5682 - val\_accuracy: 0.2750

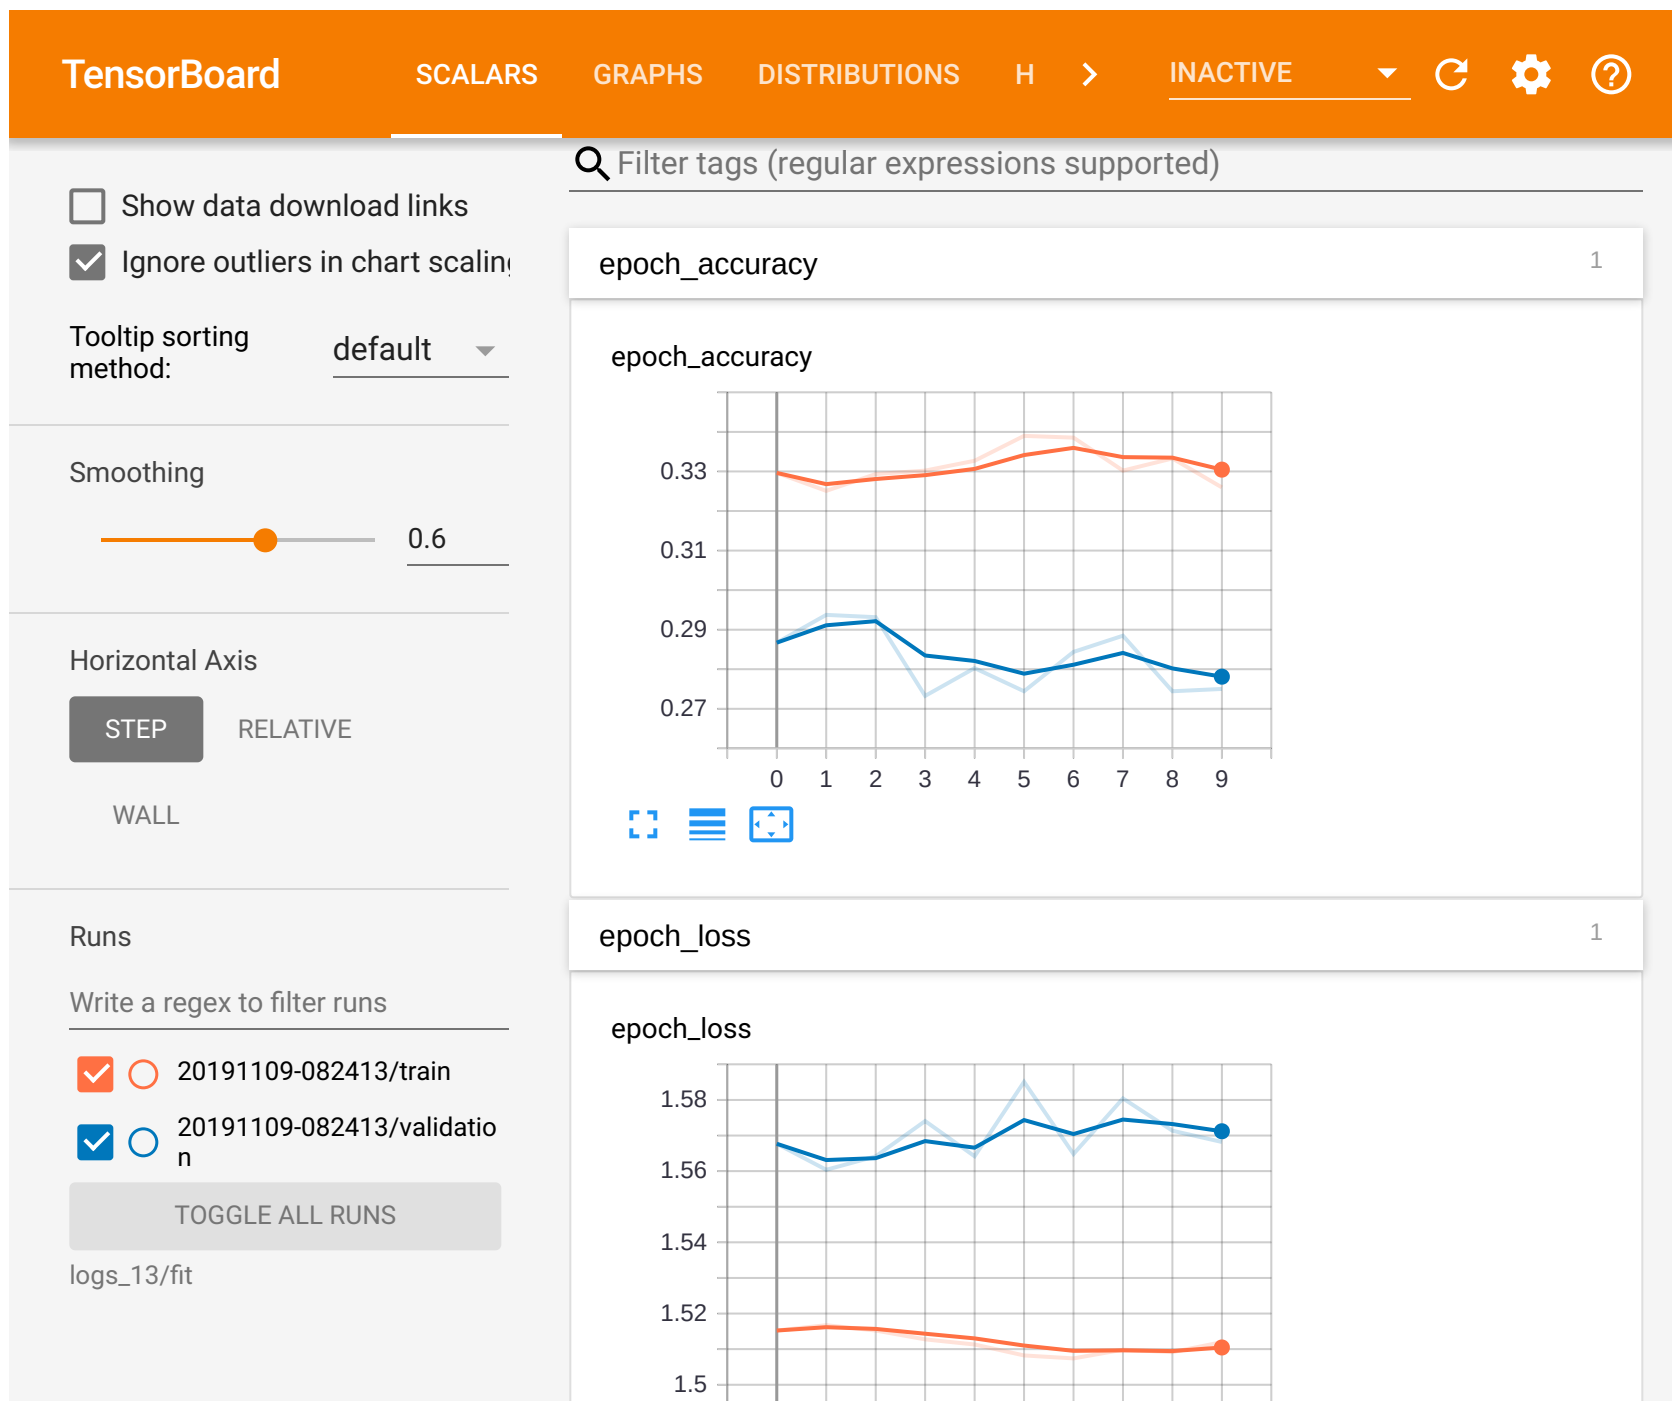
Out[166]: <tensorflow.python.keras.callbacks.History at 0x7f3eef9b9910>

In [ ]:

In [167]: *# Loads the weights*  
model\_bid.load\_weights(checkpoint\_path)  
  
loss, acc = model\_bid.evaluate(test\_X\_t, test\_y, verbose=2)  
print("Test set Accuracy: {:.2f}%".format(100\*acc))  
  
2210/1 - 2s - loss: 1.7128 - accuracy: 0.2986  
Test set Accuracy: 29.86%



```
In [168]: %tensorboard --logdir logs_13/fit
```





In [ ]:

In [ ]:

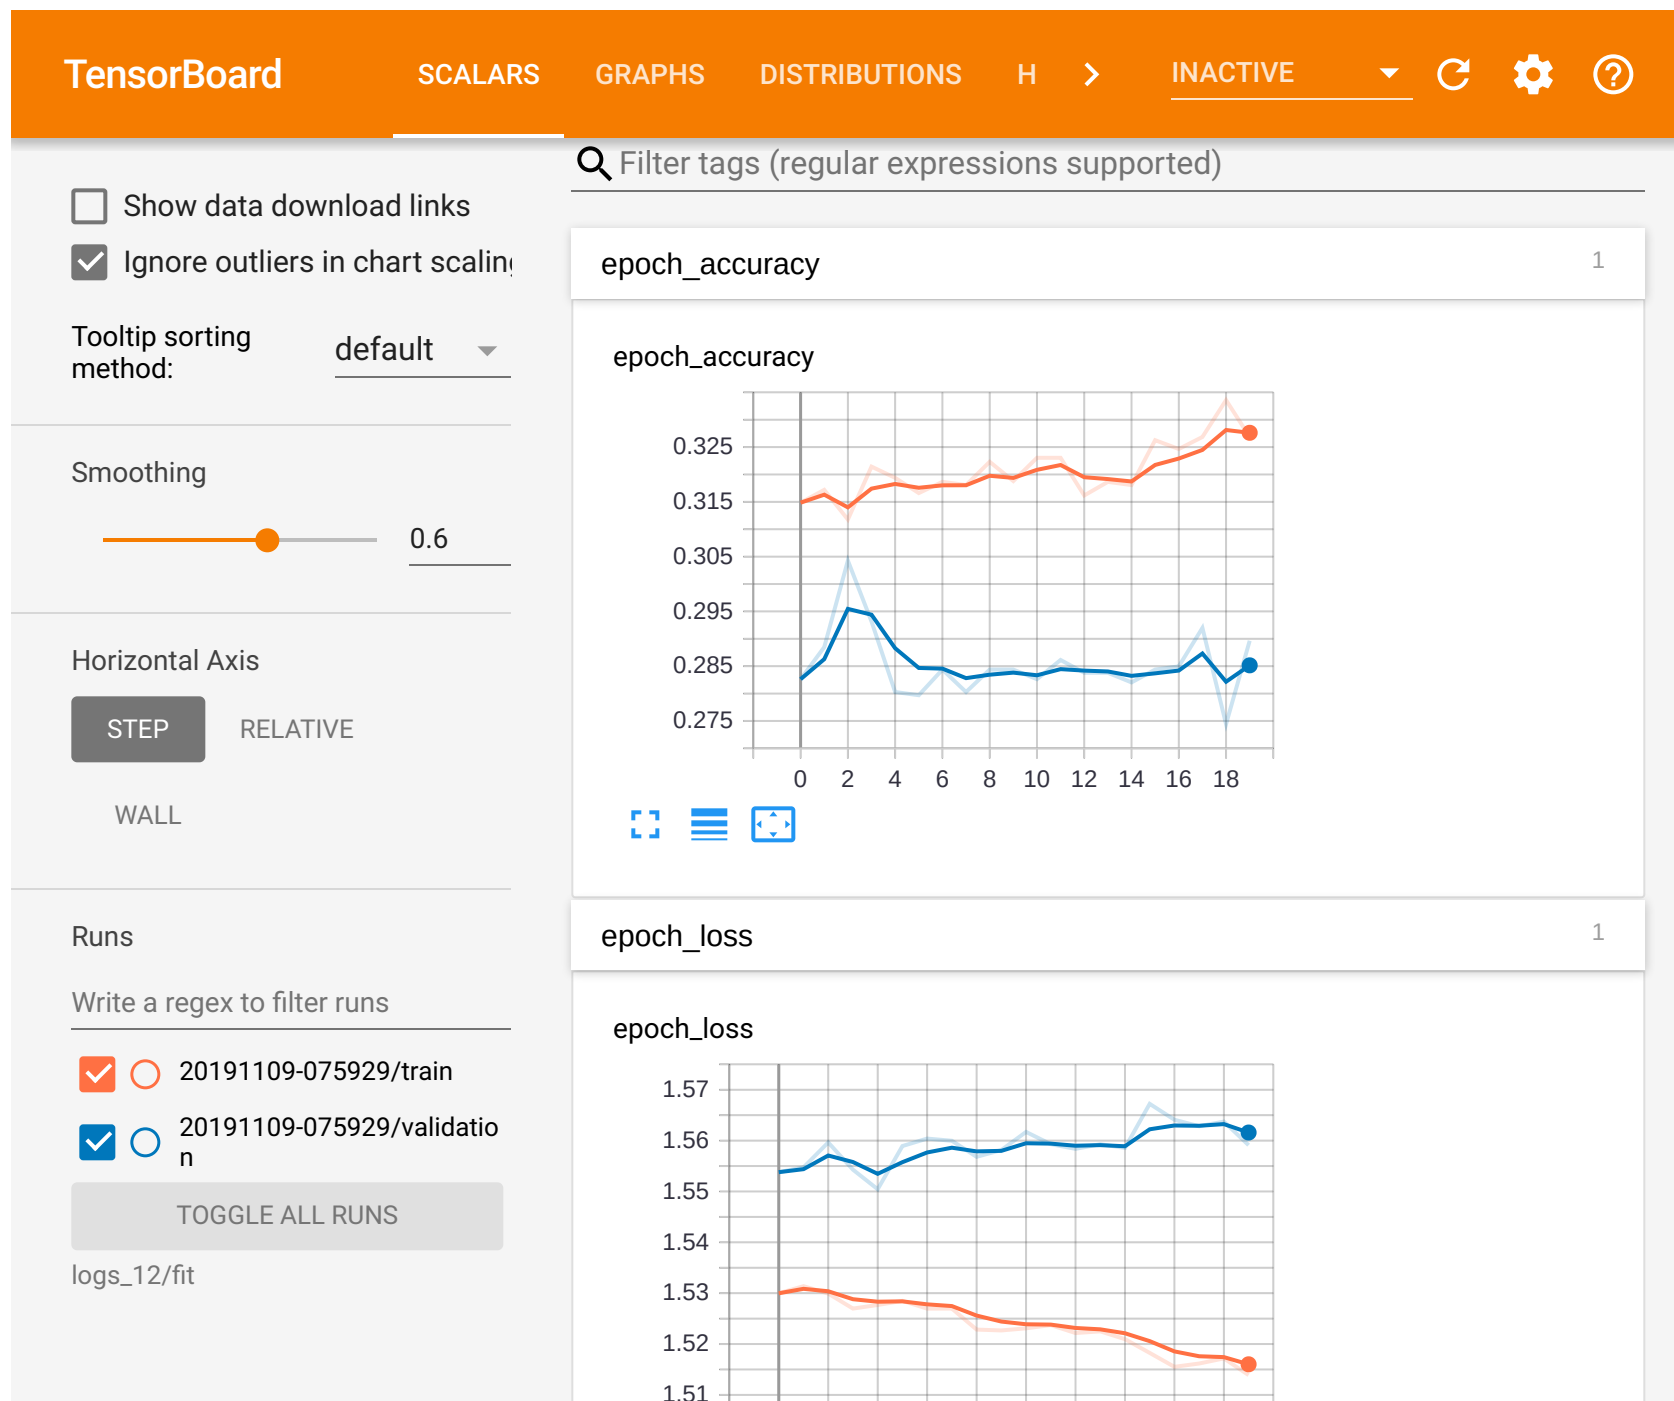
### evaluation on test data


```
In [ ]: # Loads the weights
model_bid.load_weights(checkpoint_path)

loss, acc = model_bid.evaluate(test_X_t, test_y, verbose=2)
print("Test set Accuracy: {:.2f}%".format(100*acc))
```

**Test set Accuracy: 29.86% - Using Bidirectional LSTM**

```
In [165]: %tensorboard --logdir logs_12/fit
```



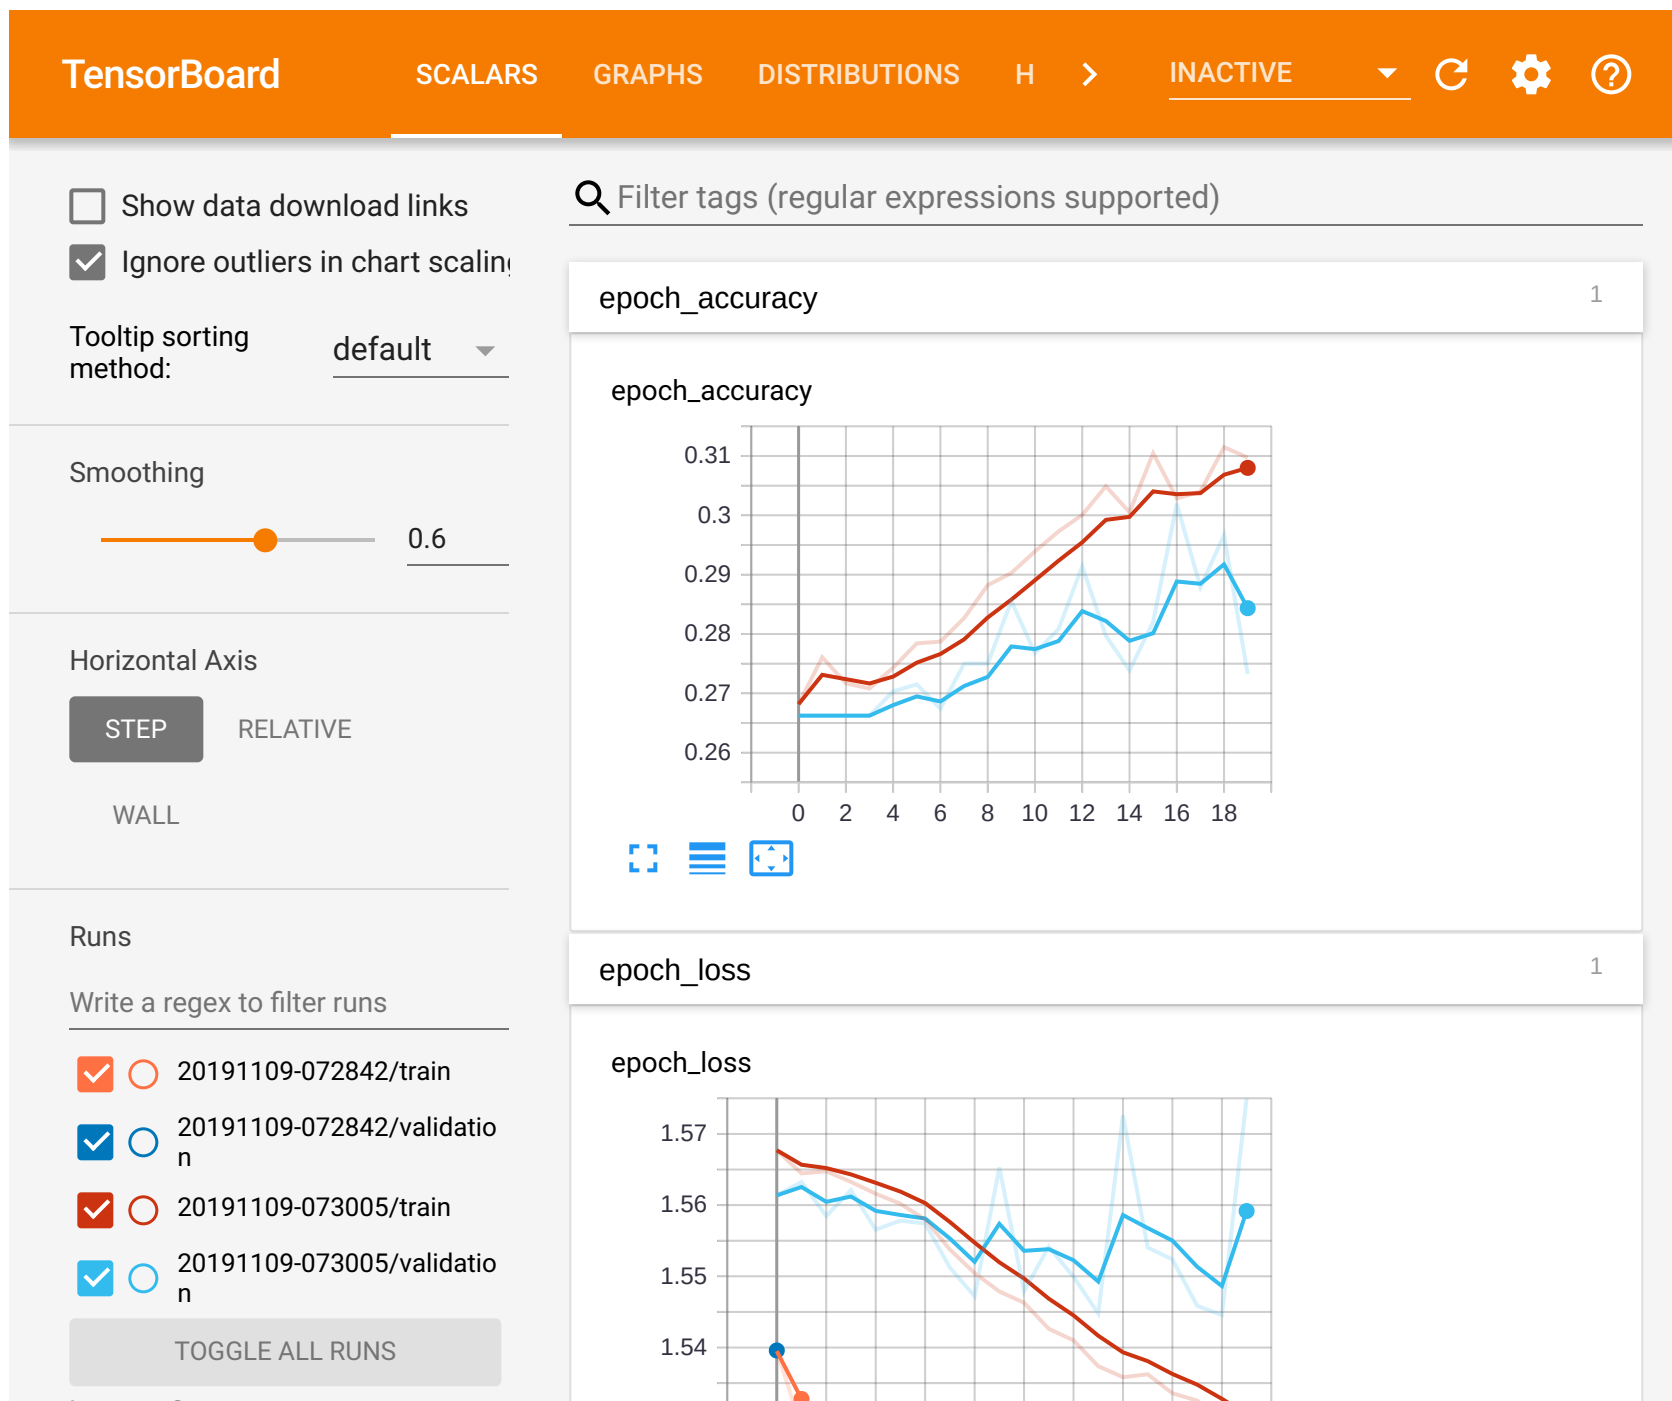



In [ ]:

In [ ]:

**Bidirectional LSTM with 20 epochs**

```
In [160]: %tensorboard --logdir logs_11/fit
```





In [ ]:

In [ ]:

### Bidirectional - with 10 epochs

In [178]: `%tensorboard --logdir logs_10/fit`

ERROR: Failed to launch TensorBoard (exited with 255).

Contents of stderr:

E1109 08:42:54.047511 140117797828416 program.py:226] TensorBoard could not bind to any port around 6006 (tried 10 times)

ERROR: TensorBoard could not bind to any port around 6006 (tried 10 times)

In [ ]:

In [ ]:

### Task 4 - Evaluate best model on test data

In [ ]:

In [ ]: `score, acc = model_lstm.evaluate(test_X_t, test_y)`In [358]: `acc`

Out[358]: 0.28597286

**Accuracy of the best model on test data = 28.60%**

In [ ]:

In [ ]:

[illegible]

In [ ]:





```
In [230]: predictions_lstm
```

```
Out[230]: array([1, 3, 1, ..., 3, 3, 1])
```

```
In [379]: predictions_lstm[-10:]
```

```
Out[379]: array([1, 1, 1, 1, 3, 3, 4, 4, 1, 1])
```

```
In [368]: len(test_y)
```

```
Out[368]: 2210
```

```
In [233]: cm_lstm = metrics.confusion_matrix(test['truth'].values, predictions_lstm)
          print(cm_lstm)
```

```
[[ 0  0  0  0  0  0]
 [ 1 93  6 173  6  0]
 [ 0 177 14 432 10  0]
 [ 0 114 11 259  5  0]
 [ 0 117  6 374 13  0]
 [ 2  81  5 302  9  0]]
```

```
In [ ]:
```

```
In [ ]:
```

```
In [234]: from sklearn.metrics import classification_report
```

```
In [235]: print(classification_report(test['truth'].values, predictions_lstm))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.16	0.33	0.22	279
2	0.33	0.02	0.04	633
3	0.17	0.67	0.27	389
4	0.30	0.03	0.05	510
5	0.00	0.00	0.00	399
accuracy			0.17	2210
macro avg	0.16	0.17	0.10	2210
weighted avg	0.22	0.17	0.10	2210

```
/home/demolakstate/anaconda3/envs/nlp/lib/python3.7/site-packages/sklearn/metrics/classification.p
y:1437: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in label
s with no predicted samples.
  'precision', 'predicted', average, warn_for)
/home/demolakstate/anaconda3/envs/nlp/lib/python3.7/site-packages/sklearn/metrics/classification.p
y:1439: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels w
ith no true samples.
  'recall', 'true', average, warn_for)
```

```
In [ ]:
```

## Evaluation of best GRU model on the test set

```
In [249]: predictions_gru = np.argmax(model_gru.predict(test_X_t), axis=1)

predictions_gru

cm_gru = metrics.confusion_matrix(test['truth'].values, predictions_gru)
print(cm_gru)

print(classification_report(test['truth'].values, predictions_gru))
```

```
[[ 0  0  0  0  0  0]
 [ 37 88 53 76 25 0]
 [ 79 198 126 181 49 0]
 [ 38 88 87 142 34 0]
 [ 38 87 89 205 91 0]
 [ 20 45 55 183 96 0]]
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.17	0.32	0.22	279
2	0.31	0.20	0.24	633
3	0.18	0.37	0.24	389
4	0.31	0.18	0.23	510
5	0.00	0.00	0.00	399
accuracy			0.20	2210
macro avg	0.16	0.18	0.16	2210
weighted avg	0.21	0.20	0.19	2210

```
/home/demolakstate/anaconda3/envs/nlp/lib/python3.7/site-packages/sklearn/metrics/classification.p
y:1437: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in label
s with no predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
/home/demolakstate/anaconda3/envs/nlp/lib/python3.7/site-packages/sklearn/metrics/classification.p
y:1439: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels w
ith no true samples.
```

```
'recall', 'true', average, warn_for)
```

```
In [ ]:
```

## Evaluation of best RNN model on the test set

```
In [255]: predictions_rnn = np.argmax(model_rnn_3.predict(test_X_t), axis=1)

predictions_rnn

cm_rnn = metrics.confusion_matrix(test['truth'].values, predictions_rnn)
print(cm_rnn)

print(classification_report(test['truth'].values, predictions_rnn))
```

```
[[ 9  2 184  84  0]
 [ 12  6 431 184  0]
 [ 17  7 268  97  0]
 [  6  3 349 152  0]
 [  0  4 282 113  0]]
      precision    recall  f1-score   support

     1         0.20      0.03      0.06         279
     2         0.27      0.01      0.02         633
     3         0.18      0.69      0.28         389
     4         0.24      0.30      0.27         510
     5         0.00      0.00      0.00         399

 accuracy          0.20         2210
 macro avg         0.18         0.21      0.12         2210
 weighted avg         0.19         0.20      0.12         2210
```

```
/home/demolakstate/anaconda3/envs/nlp/lib/python3.7/site-packages/sklearn/metrics/classification.p
y:1437: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in label
s with no predicted samples.
'precision', 'predicted', average, warn_for)
```

In [ ]:

## Discussion of the different models performance

**Bidirectional LSTM had the highest accuracy of 29.86% on the test data. This is followed by LSTM**

and RNN both with 28.60% and finally GRU model of 27.69%. MODEL\_RNN\_3 did well in predicting sentences of label 3 - "Neutral" with an F-Score of 28%. The model did badly on all other classes, and most especially on label 5. MODEL\_GRU did well on labels 3 and 1 with F-Score of 24% and 22% respectively, but badly on the others and especially on label 5. MODEL\_LSTM did well on label 3, but badly on all others and especially label 5. In conclusion, all the experimented model architectures were able to do a good job in classifying sentence sentiments that are "Neutral". They all struggled with "Strongly Positive" sentiments.

In [ ]:

In [ ]:

## Task 4 - Text Generation Language Model

In [ ]:

In [ ]:

```
In [183]: from keras.callbacks import ModelCheckpoint  
         from keras.utils import np_utils
```

In [ ]:

In [ ]:

```
In [184]: raw_text = df['text'].str.lower()
```

```
In [185]: raw_text
```

```
Out[185]: 0      the rock is destined to be the 21st century 's...
          1      the gorgeously elaborate continuation of `` th...
          2      singer/composer bryan adams contributes a slew...
          3      you 'd think by now america would have had eno...
          4      yet the act is still charming here .

          ...
          8539      a real snooze .
          8540      no surprises .
          8541      we 've seen the hippie-turned-yuppie plot befo...
          8542      her fans walked out muttering words like `` ho...
          8543      in this case zero .
Name: text, Length: 8544, dtype: object
```

```
In [186]: # create mapping of unique chars to integers
chars = sorted(list(set(raw_text)))

char_to_int = dict((c, i) for i, c in enumerate(chars))
```

```
In [187]: # summarize the loaded data
n_chars = len(raw_text)
n_vocab = len(chars)

print("Total Characters: ", n_chars)
print("Total Vocab: ", n_vocab)
```

```
Total Characters: 8544
Total Vocab: 8534
```

```
In [188]: # prepare the dataset of input to output pairs encoded as integers
seq_length = 100
dataX = []
dataY = []
for i in range(0, n_chars - seq_length, 1):
    seq_in = raw_text[i:i + seq_length]
    seq_out = raw_text[i + seq_length]
    dataX.append([char_to_int[char] for char in seq_in])
    dataY.append(char_to_int[seq_out])

n_patterns = len(dataX)
print("Total Patterns: ", n_patterns)
```

Total Patterns: 8444

```
In [189]: # reshape X to be [samples, time steps, features]
X = np.reshape(dataX, (n_patterns, seq_length, 1))

# normalize
X = X / float(n_vocab)
# one hot encode the output variable
y = np_utils.to_categorical(dataY)
```

In [ ]:

```
In [190]: # define the LSTM model
model = Sequential()
model.add(LSTM(256, input_shape=(X.shape[1], X.shape[2]), return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(256))
model.add(Dropout(0.2))
model.add(Dense(y.shape[1], activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam')
```

In [ ]:

```
In [436]: # define the checkpoint
filepath="Weights-LSTM-improvement-{epoch:02d}-{loss:.4f}-bigger.hdfs"
checkpoint = ModelCheckpoint(filepath, monitor='loss', verbose=1, save_best_only=True, mode='min')

callbacks_list = [checkpoint]

# fit the model

model.fit(X, y, epochs=80, batch_size=64, callbacks=callbacks_list)
```

```
8444/8444 [=====] - 52s 6ms/step - loss: 9.0822
```

```
Epoch 00002: loss did not improve from 9.06460
```

```
Epoch 3/80
```

```
8444/8444 [=====] - 52s 6ms/step - loss: 9.0538
```

```
Epoch 00003: loss improved from 9.06460 to 9.05378, saving model to Weights-LSTM-improvement-03-9.0538-bigger.hdfs
```

```
Epoch 4/80
```

```
8444/8444 [=====] - 52s 6ms/step - loss: 9.0529
```

```
Epoch 00004: loss improved from 9.05378 to 9.05290, saving model to Weights-LSTM-improvement-04-9.0529-bigger.hdfs
```

```
Epoch 5/80
```

```
8444/8444 [=====] - 52s 6ms/step - loss: 9.0522
```

```
Epoch 00005: loss improved from 9.05290 to 9.05219, saving model to Weights-LSTM-improvement-05-9.0522-bigger.hdfs
```

```
Epoch 6/80
```

```
8444/8444 [=====] - 52s 6ms/step - loss: 9.0516
```

```
In [ ]:
```



```
In [438]: # Load Larger LSTM network and generate text

import sys

int_to_char = dict((i, c) for i, c in enumerate(chars))
# summarize the loaded data
n_chars = len(raw_text)
n_vocab = len(chars)
print("Total Characters: ", n_chars)
print("Total Vocab: ", n_vocab)

# prepare the dataset of input to output pairs encoded as integers
seq_length = 100
dataX = []
dataY = []

for i in range(0, n_chars - seq_length, 1):
    seq_in = raw_text[i:i + seq_length]
    seq_out = raw_text[i + seq_length]
    dataX.append([char_to_int[char] for char in seq_in])
    dataY.append(char_to_int[seq_out])

n_patterns = len(dataX)
print("Total Patterns: ", n_patterns)
# reshape X to be [samples, time steps, features]
X = np.reshape(dataX, (n_patterns, seq_length, 1))

# normalize
X = X / float(n_vocab)
# one hot encode the output variable
y = np_utils.to_categorical(dataY)

# define the LSTM model
model = Sequential()
model.add(LSTM(256, input_shape=(X.shape[1], X.shape[2]), return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(256))
model.add(Dropout(0.2))
model.add(Dense(y.shape[1], activation='softmax'))

# Load the network weights
```

```

filename = "Weights-LSTM-improvement-80-9.0447-bigger.hdfs"
#filename = "Weights-LSTM-improvement-44-1.3592-bigger.hdfs"
model.load_weights(filename)

model.compile(loss='categorical_crossentropy', optimizer='adam')

# pick a random seed

start = np.random.randint(0, len(dataX)-1)
pattern = dataX[start]
print("Seed:")
print("\n", ''.join([int_to_char[value] for value in pattern]), "\n")
# generate characters
for i in range(700):
    x = np.reshape(pattern, (1, len(pattern), 1))
    x = x / float(n_vocab)
    prediction = model.predict(x, verbose=0)
    index = np.argmax(prediction)
    result = int_to_char[index]
    seq_in = [int_to_char[value] for value in pattern]
    sys.stdout.write(result)
    pattern.append(index)
    pattern = pattern[1:len(pattern)]

print("\nDone")

```

Total Characters: 8544

Total Vocab: 8534

Total Patterns: 8444

Seed:

" a generic family comedy unlikely to be appreciated by anyone outside the under-10 set .kung pow seems like some futile concoction that was developed hastily after oedekerk and his fellow moviemakers got through crashing a college keg party .kurys seems intimidated by both her subject matter and the period trappings of this debut venture into the heritage business .the film virtually chokes on its own self-consciousness .a manipulative feminist empowerment tale thinly posing as a serious drama about spousal abuse .everything in maid in manhattan is exceedingly pleasant , designed not to offend .it goes down easy , leaving virtually no aftertaste .a profoundly stupid affair , populating its hackneyed and meanspirited storyline with cardboard characters and performers who value cash above credibility .... pays tribute to heroes the way julia roberts hands out awards -- with phony humility barely camouflaging grotesque narcissism .time stands still in more ways than one in clockstoppers , a sci-fi thriller as lazy as it is interminable .as a director , eastwood is off his game -- there 's no real sense of suspense , and none of the plot ` surprises ' are really surprising .eccentric enough to stave off doldrums , caruso 's self-conscious debut is

also eminently forgettable .to work , love stories require the full emotional involvement and support of a viewer .that is made almost impossible by events that set the plot in motion .although

In [ ]:

In [ ]:

In [ ]:

**100 epochs**

In [ ]:

```

In [57]: # define the checkpoint
filepath="Weights-LSTM-improvement-{epoch:02d}-{loss:.4f}-bigger.hdfs"
checkpoint = ModelCheckpoint(filepath, monitor='loss', verbose=1, save_best_only=True, mode='min')

callbacks_list = [checkpoint]

# fit the model

model.fit(X, y, epochs=100, batch_size=64, callbacks=callbacks_list)
Epoch 00093: loss improved from 0.01940 to 0.01827, saving model to Weights-LSTM-improvement-93-0.0183-bigger.hdfs
Epoch 94/100
8444/8444 [=====] - 52s 6ms/step - loss: 0.0185

Epoch 00094: loss did not improve from 0.01827
Epoch 95/100
8444/8444 [=====] - 52s 6ms/step - loss: 0.0175

Epoch 00095: loss improved from 0.01827 to 0.01750, saving model to Weights-LSTM-improvement-95-0.0175-bigger.hdfs
Epoch 96/100
8444/8444 [=====] - 52s 6ms/step - loss: 0.0165

Epoch 00096: loss improved from 0.01750 to 0.01653, saving model to Weights-LSTM-improvement-96-0.0165-bigger.hdfs
Epoch 97/100
8444/8444 [=====] - 52s 6ms/step - loss: 0.0152

Epoch 00097: loss improved from 0.01653 to 0.01520, saving model to Weights-LSTM-improvement-97-

```

In [ ]:

```
In [60]: # Load Larger LSTM network and generate text

import sys

int_to_char = dict((i, c) for i, c in enumerate(chars))
# summarize the loaded data
n_chars = len(raw_text)
n_vocab = len(chars)
print("Total Characters: ", n_chars)
print("Total Vocab: ", n_vocab)

# prepare the dataset of input to output pairs encoded as integers
seq_length = 100
dataX = []
dataY = []

for i in range(0, n_chars - seq_length, 1):
    seq_in = raw_text[i:i + seq_length]
    seq_out = raw_text[i + seq_length]
    dataX.append([char_to_int[char] for char in seq_in])
    dataY.append(char_to_int[seq_out])

n_patterns = len(dataX)
print("Total Patterns: ", n_patterns)
# reshape X to be [samples, time steps, features]
X = np.reshape(dataX, (n_patterns, seq_length, 1))

# normalize
X = X / float(n_vocab)
# one hot encode the output variable
y = np_utils.to_categorical(dataY)

# define the LSTM model
model = Sequential()
model.add(LSTM(256, input_shape=(X.shape[1], X.shape[2]), return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(256))
model.add(Dropout(0.2))
model.add(Dense(y.shape[1], activation='softmax'))

# Load the network weights
```

```

filename = "Weights-LSTM-improvement-80-9.0447-bigger.hdfs"
#filename = "Weights-LSTM-improvement-44-1.3592-bigger.hdfs"
model.load_weights(filename)

model.compile(loss='categorical_crossentropy', optimizer='adam')

# pick a random seed

start = np.random.randint(0, len(dataX)-1)
pattern = dataX[start]
print("Seed:")
print("\n", ''.join([int_to_char[value] for value in pattern]), "\n")
# generate characters
for i in range(70):
    x = np.reshape(pattern, (1, len(pattern), 1))
    x = x / float(n_vocab)
    prediction = model.predict(x, verbose=0)
    index = np.argmax(prediction)
    result = int_to_char[index]
    seq_in = [int_to_char[value] for value in pattern]
    sys.stdout.write(result)
    pattern.append(index)
    pattern = pattern[1:len(pattern)]

print("\nDone")

```

he line , to find a place among the studio 's animated classics .slow and ponderous , but rohmer 's drama builds to an intense indoor drama about compassion , sacrifice , and christian love in t he face of political corruption .if you 're not totally weirded - out by the notion of cinema as community-therapy spectacle , quitting hits home with disorienting force .austin powers for the m ost part is extremely funny , the first part making up for any flaws that come later .while tatto o borrows heavily from both seven and the silence of the lambs , it manages to maintain both a le vel of sophisticated intrigue and human-scale characters that suck the audience in .cho continues her exploration of the outer limits of raunch with considerable brio .elvira fans could hardly as k for more .a canny , derivative , wildly gruesome portrait of a london sociopath who 's the scar iest of sadists .the movie should be credited with remembering his victims .fast-paced and wonder fully edited , the film is extremely thorough .a bracing , unblinking work that serves as a painf ul elegy and sobering cautionary tale .hashiguchi uses the situation to evoke a japan bustling at op an undercurrent of loneliness and isolation .as if trying to grab a lump of play-doh , the har der that liman tries to squeeze his story , the more details slip out between his fingers .my big fat greek wedding is not only the best date movie of the year , it 's also a -- dare i say it twi ce -- delightfully charming -- and totally american , i might add -- slice of comedic bliss .few films have captured the chaos of an urban conflagration with such fury , and audience members wil l leave feeling as shaken as nesbitt 's cooper looks when the bullets stop flying .another love s

tory in 2002 's remarkable procession of sweeping pictures that have reinvigorated the romance ge

In [ ]:

In [ ]:

## 500 epochs

In [ ]:

```
In [61]: # define the LSTM model
model_500 = Sequential()
model_500.add(LSTM(256, input_shape=(X.shape[1], X.shape[2]), return_sequences=True))
model_500.add(Dropout(0.2))
model_500.add(LSTM(256))
model_500.add(Dropout(0.2))
model_500.add(Dense(y.shape[1], activation='softmax'))
model_500.compile(loss='categorical_crossentropy', optimizer='adam')
```

In [ ]:

In [ ]:

```
In [62]: # define the checkpoint
filepath="Weights-LSTM-improvement-{epoch:02d}-{loss:.4f}-bigger.hdfs"
checkpoint = ModelCheckpoint(filepath, monitor='loss', verbose=1, save_best_only=True, mode='min')

callbacks_list = [checkpoint]

# fit the model

model_500.fit(X, y, epochs=500, batch_size=64, callbacks=callbacks_list)
```

```
Epoch 00496: loss did not improve from 9.02926
Epoch 497/500
8444/8444 [=====] - 51s 6ms/step - loss: 9.0445

Epoch 00497: loss did not improve from 9.02926
Epoch 498/500
8444/8444 [=====] - 51s 6ms/step - loss: 9.0445

Epoch 00498: loss did not improve from 9.02926
Epoch 499/500
8444/8444 [=====] - 51s 6ms/step - loss: 9.0445

Epoch 00499: loss did not improve from 9.02926
Epoch 500/500
8444/8444 [=====] - 51s 6ms/step - loss: 9.0445

Epoch 00500: loss did not improve from 9.02926
```

```
Out[62]: <keras.callbacks.callbacks.History at 0x7f3df1541e90>
```

```
In [ ]:
```



```
In [63]: # Load Larger LSTM network and generate text

import sys

int_to_char = dict((i, c) for i, c in enumerate(chars))
# summarize the loaded data
n_chars = len(raw_text)
n_vocab = len(chars)
print("Total Characters: ", n_chars)
print("Total Vocab: ", n_vocab)

# prepare the dataset of input to output pairs encoded as integers
seq_length = 100
dataX = []
dataY = []

for i in range(0, n_chars - seq_length, 1):
    seq_in = raw_text[i:i + seq_length]
    seq_out = raw_text[i + seq_length]
    dataX.append([char_to_int[char] for char in seq_in])
    dataY.append(char_to_int[seq_out])

n_patterns = len(dataX)
print("Total Patterns: ", n_patterns)
# reshape X to be [samples, time steps, features]
X = np.reshape(dataX, (n_patterns, seq_length, 1))

# normalize
X = X / float(n_vocab)
# one hot encode the output variable
y = np_utils.to_categorical(dataY)

# define the LSTM model
model_500 = Sequential()
model_500.add(LSTM(256, input_shape=(X.shape[1], X.shape[2]), return_sequences=True))
model_500.add(Dropout(0.2))
model_500.add(LSTM(256))
model_500.add(Dropout(0.2))
model_500.add(Dense(y.shape[1], activation='softmax'))

# Load the network weights
```



100 movies of the year. Name of the worst movies of the year. ?

Done

In [ ]:

In [ ]:

In [ ]:

## 1000 epochs

In [ ]:

```
In [191]: # define the LSTM model
model_1000 = Sequential()
model_1000.add(LSTM(256, input_shape=(X.shape[1], X.shape[2]), return_sequences=True))
model_1000.add(Dropout(0.2))
model_1000.add(LSTM(256))
model_1000.add(Dropout(0.2))
model_1000.add(Dense(y.shape[1], activation='softmax'))
model_1000.compile(loss='categorical_crossentropy', optimizer='adam')
```

In [ ]:

In [ ]:

```
In [192]: # define the checkpoint
filepath="Weights-LSTM-improvement-{epoch:02d}-{loss:.4f}-bigger.hdfs"
checkpoint = ModelCheckpoint(filepath, monitor='loss', verbose=1, save_best_only=True, mode='min')

callbacks_list = [checkpoint]

# fit the model

model_1000.fit(X, y, epochs=1000, batch_size=64, callbacks=callbacks_list)
```

```
Epoch 00996: loss did not improve from 9.00526
Epoch 997/1000
8444/8444 [=====] - 52s 6ms/step - loss: 9.0445

Epoch 00997: loss did not improve from 9.00526
Epoch 998/1000
8444/8444 [=====] - 52s 6ms/step - loss: 9.0445

Epoch 00998: loss did not improve from 9.00526
Epoch 999/1000
8444/8444 [=====] - 52s 6ms/step - loss: 9.0445

Epoch 00999: loss did not improve from 9.00526
Epoch 1000/1000
8444/8444 [=====] - 52s 6ms/step - loss: 9.0445

Epoch 01000: loss did not improve from 9.00526
```

```
Out[192]: <keras.callbacks.callbacks.History at 0x7f3f26f634d0>
```

```
In [ ]:
```

```
In [193]: # Load Larger LSTM network and generate text

import sys

int_to_char = dict((i, c) for i, c in enumerate(chars))
# summarize the loaded data
n_chars = len(raw_text)
n_vocab = len(chars)
print("Total Characters: ", n_chars)
print("Total Vocab: ", n_vocab)

# prepare the dataset of input to output pairs encoded as integers
seq_length = 100
dataX = []
dataY = []

for i in range(0, n_chars - seq_length, 1):
    seq_in = raw_text[i:i + seq_length]
    seq_out = raw_text[i + seq_length]
    dataX.append([char_to_int[char] for char in seq_in])
    dataY.append(char_to_int[seq_out])

n_patterns = len(dataX)
print("Total Patterns: ", n_patterns)
# reshape X to be [samples, time steps, features]
X = np.reshape(dataX, (n_patterns, seq_length, 1))

# normalize
X = X / float(n_vocab)
# one hot encode the output variable
y = np_utils.to_categorical(dataY)

# define the LSTM model
model_1000 = Sequential()
model_1000.add(LSTM(256, input_shape=(X.shape[1], X.shape[2]), return_sequences=True))
model_1000.add(Dropout(0.2))
model_1000.add(LSTM(256))
model_1000.add(Dropout(0.2))
model_1000.add(Dense(y.shape[1], activation='softmax'))

# Load the network weights
```

```

filename = "Weights-LSTM-improvement-80-9.0447-bigger.hdfs"
#filename = "Weights-LSTM-improvement-44-1.3592-bigger.hdfs"
model_1000.load_weights(filename)

model_1000.compile(loss='categorical_crossentropy', optimizer='adam')

# pick a random seed

start = np.random.randint(0, len(dataX)-1)
pattern = dataX[start]
print("Seed:")
print("\n", ''.join([int_to_char[value] for value in pattern]), "\n")
# generate characters
for i in range(70):
    x = np.reshape(pattern, (1, len(pattern), 1))
    x = x / float(n_vocab)
    prediction = model_1000.predict(x, verbose=0)
    index = np.argmax(prediction)
    result = int_to_char[index]
    seq_in = [int_to_char[value] for value in pattern]
    sys.stdout.write(result)
    pattern.append(index)
    pattern = pattern[1:len(pattern)]

print("\nDone")

```

s a self-aware , often self-mocking , intelligence .the chateau is a risky venture that never quite goes where you expect and often surprises you with unexpected comedy .a very well-meaning movie , and it will stand in future years as an eloquent memorial to the world trade center tragedy . there are n't many conclusive answers in the film , but there is an interesting story of pointed personalities , courage , tragedy and the little guys vs. the big guys .vividly demonstrates that the director of such hollywood blockbusters as patriot games can still turn out a small , personal film with an emotional wallop .a four star performance from kevin kline who unfortunately works with a two star script .dogtown & z-boys evokes the blithe rebel fantasy with the kind of insouciance embedded in the sexy demise of james dean .if you do n't flee , you might be seduced .if you do n't laugh , flee .payne constructs a hilarious ode to middle america and middle age with this unlikely odyssey , featuring a pathetic , endearing hero who is all too human .koury frighteningly and honestly exposes one teenager 's uncomfortable class resentment and , in turn , his self-inflicted retaliation .the santa clause 2 proves itself a more streamlined and thought out encounter than the original could ever have hoped to be .now as a former gong show addict , i 'll admit it , my only complaint is that we did n't get more re-creations of all those famous moments from the show .succeeds where its recent predecessor miserably fails because it demands that you suffer the dreadfulness of war from both sides .the first bond movie in ages that is n't fake fun .this odd , poetic road movie , spiked by jolts of pop music , pretty much takes place in morton 's even watchful gaze , and it is a tribute to the actress , and to her inventive director , that the

r-watchful gaze -- and it is a tribute to the actress, and to her inventive director, that the journey is such a mesmerizing one. a film centering on a traditional indian wedding in contempora

In [ ]:

In [ ]:

## Qualitative Discussion of the model

**All the models did pretty well at the initial text character generation. The model trained with 100 epochs, unlike that of 500 epochs and 1000 epochs, did not make so much sense at the end of the first paragraph. All the models failed at the generation of the second paragraph.**

In [ ]: