

gglImage Disaster Images

In []:

In []:

Discussion

Looking at the heatmap, the model does not perform as well on 1s as on 0s considering on the diagonal, the 1s is darker. On the error plot, the column for class 0 is quite bright, which tells that many images get missclassified as 0s. The model performs less when only the last layer is fine-tuned with test set accuracy of 79.201% compared to 81.531% when more layers are fine-tuned.

In []:

The dataset size is not a representative of how well the model performs. The highest test-set accuracies are not necessarily the largest dataset size

In []:

In []:

In []:

In [42]:

```
from __future__ import absolute_import, division, print_function, unicode_literals
```

```
In [43]: import tensorflow as tf  
tf.test.gpu_device_name()
```

Out[43]: ''

```
In [44]: from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Conv2D, Flatten, Dropout, MaxPooling2D  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
  
import os  
import numpy as np  
import matplotlib.pyplot as plt
```

In []:

In []:

```
In [45]: import os
```

```
In [46]: from sklearn.metrics import confusion_matrix
```

In []:

In []:

In []:

Load data

```
In [47]: train_dir = 'data_ggImage/train'  
validation_dir = 'data_ggImage/validation'  
test_dir = 'data_ggImage/test'
```

In []:

```
In [48]: train_damaged_dir = os.path.join(train_dir, 'damaged') # directory with our training damaged pictures  
train_undamaged_dir = os.path.join(train_dir, 'undamaged') # directory with our training undamaged pictures  
validation_damaged_dir = os.path.join(validation_dir, 'damaged') # directory with our validation damaged pictures  
validation_undamaged_dir = os.path.join(validation_dir, 'undamaged') # directory with our validation undamaged pictures  
test_damaged_dir = os.path.join(test_dir, 'damaged') # directory with our test damaged pictures  
test_undamaged_dir = os.path.join(test_dir, 'undamaged') # directory with our test undamaged pictures
```

```
In [ ]:
```

Understand the data

Let's look at how many damaged and undamaged images are in the training and validation directory:

```
In [49]: num_damaged_tr = len(os.listdir(train_damaged_dir))  
num_undamaged_tr = len(os.listdir(train_undamaged_dir))  
  
num_damaged_val = len(os.listdir(validation_damaged_dir))  
num_undamaged_val = len(os.listdir(validation_undamaged_dir))  
  
num_damaged_ts = len(os.listdir(test_damaged_dir))  
num_undamaged_ts = len(os.listdir(test_undamaged_dir))  
  
total_train = num_damaged_tr + num_undamaged_tr  
total_val = num_damaged_val + num_undamaged_val  
total_test = num_damaged_ts + num_undamaged_ts
```

```
In [ ]:
```

```
In [50]: print('total training damaged images:', num_damaged_tr)
print('total training undamaged images:', num_undamaged_tr)

print('total validation damaged images:', num_damaged_val)
print('total validation undamaged images:', num_undamaged_val)

print('total test damaged images:', num_damaged_ts)
print('total test undamaged images:', num_undamaged_ts)
print("--")
print("Total training images:", total_train)
print("Total validation images:", total_val)
print("Total test images:", total_test)
```

```
total training damaged images: 1210
total training undamaged images: 594
total validation damaged images: 404
total validation undamaged images: 198
total test damaged images: 403
total test undamaged images: 198
--
Total training images: 1804
Total validation images: 602
Total test images: 601
```

```
In [ ]:
```

```
In [51]: # set up variables  
batch_size = 128  
#epochs = 5  
#IMG_HEIGHT = 150  
#IMG_WIDTH = 150  
  
# VGG19, VGG16, Inception  
#IMG_HEIGHT = 224  
#IMG_WIDTH = 224  
  
# inceptionV3  
#IMG_HEIGHT = 299  
#IMG_WIDTH = 299  
  
IMG_HEIGHT = 331  
IMG_WIDTH = 331
```

```
In [ ]:
```

Data preparation

```
In [52]: train_image_generator = ImageDataGenerator(rescale=1./255) # Generator for our training data  
validation_image_generator = ImageDataGenerator(rescale=1./255) # Generator for our validation data  
test_image_generator = ImageDataGenerator(rescale=1./255) # Generator for our test data
```

```
In [53]: train_data_gen = train_image_generator.flow_from_directory(batch_size=batch_size,  
                                                               directory=train_dir,  
                                                               shuffle=True,  
                                                               target_size=(IMG_HEIGHT, IMG_WIDTH),  
                                                               class_mode='binary')
```

Found 1804 images belonging to 2 classes.

Found 602 images belonging to 2 classes.

Found 601 images belonging to 2 classes.

In []:

In [1]:

Visualize training images

```
In [15]: sample training images,    = next(train data gen)
```

```
In [16]: sample_training_images, labels = next(train_data_gen)
```

In [17]: labels

```
Out[17]: array([0., 0., 0., 1., 0., 0., 0., 1., 0., 0., 0., 0., 1., 1., 1., 0., 1., 0., 0.,  
    0., 1., 0., 0., 0., 0., 1., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.,  
    0., 0., 1., 0., 0., 1., 1., 0., 1., 1., 1., 1., 1., 0., 1., 0., 0., 0., 0., 0., 0.,  
    1., 1., 0., 0., 1., 0., 1., 0., 0., 1., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 1.,  
    1., 0., 0., 0., 0., 1., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,  
    1., 1., 1., 0., 1., 0., 0., 1., 1., 1., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 1.,  
    0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 1., 0., 1., 0., 0., 0., 0., 0., 0., 1.,  
    0., 1., 0., 0., 0., 0., 0., 0., 0.], dtype=float32)
```

```
In [18]: sample training images, labels = next(train data gen)
```

In [19]: labels

In []:

In []:

```
In [20]: # This function will plot images in the form of a grid with 1 row and 5 columns where images are placed side-by-side horizontally.
def plotImages(images_arr):
    fig, axes = plt.subplots(1, 5, figsize=(20,20))
    axes = axes.flatten()
    for img, ax in zip( images_arr, axes):
        ax.imshow(img)
        ax.axis('off')
    plt.tight_layout()
    plt.show()
```

```
In [21]: plotImages(sample_training_images[:5])
```



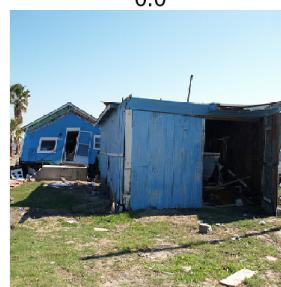
In []:

In []:

```
In [22]: def show_batch(image_batch, label_batch):
    plt.figure(figsize=(40,40))
    for n in range(25):
        ax = plt.subplot(5,5,n+1)
        plt.imshow(image_batch[n])
        #plt.title(CLASS_NAMES[label_batch[n]==1][0].title())
        plt.title(label_batch[n], fontdict={'fontsize':28})
        plt.axis('off')
```

```
In [23]: image_batch, label_batch = next(train_data_gen)
show_batch(image_batch, label_batch)
```





In []:

```
In [24]: #train_data_gen.classes[:50]
```

In []:

In [25]: label_batch

In []:

In []:

In []:

In [1]:

In [1]:

NASNetLarge Pre-trained model

```
In [92]: from keras.applications.nasnet import NASNetLarge  
from keras.preprocessing import image  
from keras.applications.nasnet import preprocess_input
```

```
In [ ]:
```

```
In [93]: from keras.models import Model  
from keras.layers import Dense, GlobalAveragePooling2D
```

```
In [ ]:
```

```
In [ ]:
```

```
In [94]: from keras.callbacks import EarlyStopping
```

```
In [ ]:
```

```
In [30]: #callback = EarlyStopping(monitor='val_loss', patience=7)  
callback = EarlyStopping(monitor='val_acc', patience=7, restore_best_weights=True)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [31]: base_model = NASNetLarge(weights='imagenet', include_top=False)
```

WARNING: Logging before flag parsing goes to stderr.

W1201 22:51:40.121912 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

W1201 22:51:40.187230 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

W1201 22:51:40.218104 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:4479: The name tf.truncated_normal is deprecated. Please use tf.random.truncated_normal instead.

W1201 22:51:40.267564 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

W1201 22:51:40.268236 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

W1201 22:51:40.519651 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:2041: The name tf.nn.fused_batch_norm is deprecated. Please use tf.compat.v1.nn.fused_batch_norm instead.

W1201 22:51:40.957061 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

W1201 22:51:41.216487 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:4271: The name tf.nn.avg_pool is deprecated. Please use tf.nn.avg_pool2d instead.

```
In [ ]:
```

```
In [32]: # add a global spatial average pooling layer
x = base_model.output
x = GlobalAveragePooling2D()(x)
# let's add a fully-connected layer
x = Dense(1024, activation='relu')(x)
# and a logistic layer -- let's say we have 200 classes
predictions = Dense(2, activation='softmax')(x)

# this is the model we will train
model_NasNet = Model(inputs=base_model.input, outputs=predictions)

# first: train only the top layers (which were randomly initialized)
# i.e. freeze all convolutional InceptionV3 layers
for layer in base_model.layers:
    layer.trainable = False

# compile the model (should be done *after* setting layers to non-trainable)
model_NasNet.compile(optimizer='rmsprop', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

W1201 22:52:19.929330 140279426041664 deprecation_wrapper.py:119] From /home/demolakstate/anaconda
3/envs/dl/lib/python3.7/site-packages/keras/optimizers.py:793: The name tf.train.Optimizer is depre
cated. Please use tf.compat.v1.train.Optimizer instead.

```
In [33]: history_NasNet = model_NasNet.fit_generator(  
    train_data_gen,  
    steps_per_epoch=total_train // batch_size,  
    epochs=15,  
    validation_data=val_data_gen,  
    validation_steps=total_val // batch_size,  
    #callbacks=[callback, cp_callback]  
    callbacks=[callback]  
)
```

W1201 22:52:20.059524 140279426041664 deprecation.py:323] From /home/demolakstate/anaconda3/envs/dl/lib/python3.7/site-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Epoch 1/15
14/14 [=====] - 757s 54s/step - loss: 2.4034 - acc: 0.6788 - val_loss: 0.8060 - val_acc: 0.7637
Epoch 2/15
14/14 [=====] - 686s 49s/step - loss: 0.7974 - acc: 0.7494 - val_loss: 0.8794 - val_acc: 0.6540
Epoch 3/15
14/14 [=====] - 675s 48s/step - loss: 0.5678 - acc: 0.7881 - val_loss: 0.8103 - val_acc: 0.7511
Epoch 4/15
14/14 [=====] - 674s 48s/step - loss: 0.6117 - acc: 0.7786 - val_loss: 0.8819 - val_acc: 0.7996
Epoch 5/15
14/14 [=====] - 675s 48s/step - loss: 0.5274 - acc: 0.8010 - val_loss: 0.6605 - val_acc: 0.7764
Epoch 6/15
14/14 [=====] - 675s 48s/step - loss: 0.4605 - acc: 0.8195 - val_loss: 1.0308 - val_acc: 0.7246
Epoch 7/15
14/14 [=====] - 720s 51s/step - loss: 0.4382 - acc: 0.8220 - val_loss: 0.8957 - val_acc: 0.7321
Epoch 8/15
14/14 [=====] - 651s 46s/step - loss: 0.5294 - acc: 0.8052 - val_loss: 0.5401 - val_acc: 0.7975
Epoch 9/15
14/14 [=====] - 698s 50s/step - loss: 0.4220 - acc: 0.8453 - val_loss: 0.5

```
031 - val_acc: 0.7890
Epoch 10/15
14/14 [=====] - 712s 51s/step - loss: 0.3597 - acc: 0.8482 - val_loss: 0.4
520 - val_acc: 0.8228
Epoch 11/15
14/14 [=====] - 712s 51s/step - loss: 0.3978 - acc: 0.8453 - val_loss: 0.5
060 - val_acc: 0.7969
Epoch 12/15
14/14 [=====] - 681s 49s/step - loss: 0.3537 - acc: 0.8509 - val_loss: 0.4
972 - val_acc: 0.7890
Epoch 13/15
14/14 [=====] - 688s 49s/step - loss: 0.3621 - acc: 0.8532 - val_loss: 0.6
312 - val_acc: 0.7468
Epoch 14/15
14/14 [=====] - 680s 49s/step - loss: 0.2092 - acc: 0.9210 - val_loss: 0.6
988 - val_acc: 0.8059
Epoch 15/15
14/14 [=====] - 695s 50s/step - loss: 0.3276 - acc: 0.8896 - val_loss: 0.5
596 - val_acc: 0.7743
```

```
In [35]: # Save the entire model to a HDF5 file
#model_NasNet.save('data_ggImage/NasNet_model_ggImage.h5')
```

```
In [ ]:
```

```
In [ ]:
```

```
In [59]: #model_NasNet.load_weights('NasNet/NasNet_model.h5')
```

```
model_NasNet = tf.keras.models.load_model('data_ggImage/NasNet_model_ggImage.h5')
```

```
W1202 18:54:34.968930 140279426041664 deprecation.py:573] From /home/demolakstate/anaconda3/envs/d  
l/lib/python3.7/site-packages/tensorflow/python/util/deprecation.py:507: calling VarianceScaling.__  
init__ (from tensorflow.python.ops.init_ops) with distribution=normal is deprecated and will be rem  
oved in a future version.
```

```
Instructions for updating:
```

```
`normal` is a deprecated alias for `truncated_normal`
```

```
W1202 18:54:34.992153 140279426041664 deprecation.py:506] From /home/demolakstate/anaconda3/envs/d  
l/lib/python3.7/site-packages/tensorflow/python/ops/init_ops.py:1251: calling VarianceScaling.__ini  
t__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future  
version.
```

```
Instructions for updating:
```

```
Call initializer instance with the dtype argument instead of passing it to the constructor
```

```
In [ ]:
```

```
In [ ]:
```

Visualize NasNetLarge model

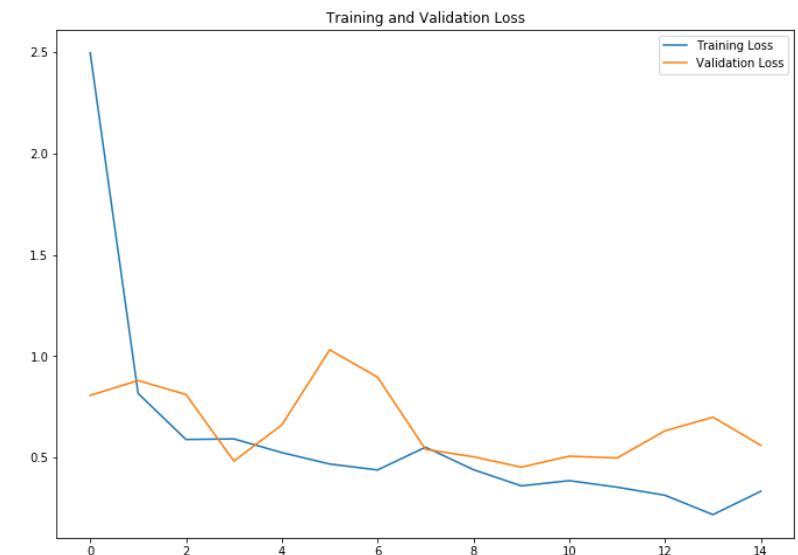
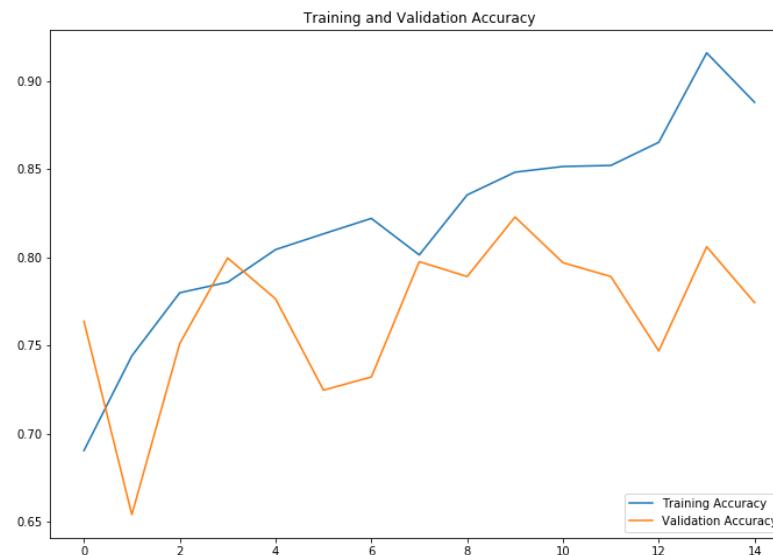
```
In [81]: acc = history_NasNet.history['acc']
val_acc = history_NasNet.history['val_acc']

loss = history_NasNet.history['loss']
val_loss = history_NasNet.history['val_loss']

epochs_range = range(15)

plt.figure(figsize=(25, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



In []:

In []:

Evaluation of NasNet model accuracy on test data

Let's compare how the model performs on the test dataset

In [60]:

```
test_loss, test_acc = model_NasNet.evaluate_generator(test_data_gen, verbose=0)
```

In [61]:

```
print('\nTest accuracy:', test_acc)
```

Test accuracy: 0.7920133

In [62]:

```
print('\nTest loss:', test_loss)
```

Test loss: 0.5111856520175934

In []:

In []:

Make predictions on test data

Let's make predictions on some images

```
In [63]: test_data_gen.class_indices
```

```
Out[63]: {'damaged': 0, 'undamaged': 1}
```

```
In [64]: predictions = model_NasNet.predict_generator(test_data_gen)
```

```
In [65]: true_labels = test_data_gen.classes
```

```
In [66]: predictions[0]
```

```
Out[66]: array([0.09723058, 0.9027694 ], dtype=float32)
```

```
In [67]: np.argmax(predictions[0])
```

```
Out[67]: 1
```

```
In [68]: test_data_gen.classes[0]
```

```
Out[68]: 0
```

```
In [69]: predictions[-1]
```

```
Out[69]: array([0.99654144, 0.00345858], dtype=float32)
```

```
In [70]: np.argmax(predictions[-1])
```

```
Out[70]: 0
```

```
In [71]: test_data_gen.classes[-1]
```

```
Out[71]: 1
```

```
In [ ]:
```

```
In [72]: test_data_gen.classes
```

In []:

In []:

Confusion Matrix

```
In [73]: import matplotlib.pyplot as plt  
import numpy as np
```

```
In [ ]:
```

```
In [ ]:
```

```
In [74]: y_true = true_labels  
y_pred = np.array([np.argmax(x) for x in predictions])
```

```
In [75]: y_pred
```

```
Out[75]: array([1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1,  
0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0,  
0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0,  
1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,  
1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,  
1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0,  
0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0,  
0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0,  
1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0,  
1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0,  
1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0,  
0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,  
1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,  
1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0,  
1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0,  
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0,  
1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1,  
1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1,  
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1,  
0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1,  
1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1,  
0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,  
0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1,  
0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1,  
0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,  
0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,  
1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0,  
0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1,  
0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,  
0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,  
0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0,
```

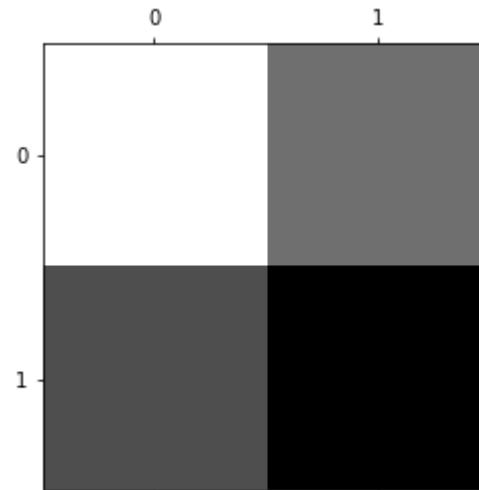
In []:

In []:

In [76]: `cm = confusion_matrix(y_true, y_pred)`In [77]: `print(cm)`

```
[[253 150]
 [127  71]]
```

In []:

In [78]: `plt.matshow(cm, cmap=plt.cm.gray)
plt.show()`

In []:

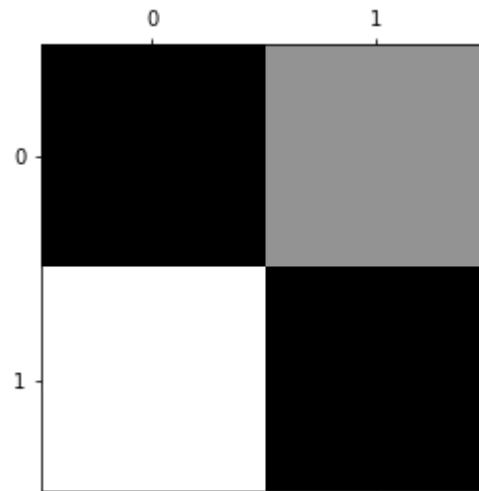
In []:

Plot on Errors

```
In [79]: row_sums = cm.sum(axis=1, keepdims=True)
norm_cm = cm / row_sums
```

```
In [ ]:
```

```
In [80]: np.fill_diagonal(norm_cm, 0)
plt.matshow(norm_cm, cmap=plt.cm.gray)
plt.show()
```



```
In [ ]:
```

```
In [ ]:
```

Confusion Matrix Heat Map

```
In [583]: #!pip install seaborn
```

```
In [584]: import seaborn as sb
```

```
In [585]: #heat_map = sb.heatmap(cm, annot=True)
#sb.set(font_scale=1)
#plt.show()
```

Classification report of NasNet on test set

```
In [586]: from sklearn.metrics import classification_report
```

```
In [587]: print(classification_report(y_true, y_pred))
```

	precision	recall	f1-score	support
0	0.61	0.62	0.62	274
1	0.42	0.41	0.41	182
accuracy			0.54	456
macro avg	0.52	0.52	0.52	456
weighted avg	0.54	0.54	0.54	456

```
In [ ]:
```

Fine-tuning more layers in NASNetLarge

```
In [82]: from keras.callbacks import EarlyStopping
```

```
In [83]: checkpoint_path = "NasNet/cp_2.ckpt"
checkpoint_dir = os.path.dirname(checkpoint_path)
```

```
In [84]: # Create a callback that save the model's weights
cp_callback = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path, save_weights_only=True, v
```

```
In [ ]:
```

```
In [ ]:
```

```
In [85]: #callback = EarlyStopping(monitor='val_loss', patience=7)

callback = EarlyStopping(monitor='val_acc', patience=5, restore_best_weights=True)
```

```
In [ ]:
```

```
In [86]: for i, layer in enumerate(base_model.layers):
    print(i, layer.name)

21 separable_conv_1_pad_reduction_right2_stem_1
22 activation_8
23 reduction_add_1_stem_1
24 separable_conv_1_reduction_right2_stem_1
25 separable_conv_1_pad_reduction_right3_stem_1
26 activation_10
27 separable_conv_1_bn_reduction_right2_stem_1
28 separable_conv_1_reduction_right3_stem_1
29 separable_conv_1_reduction_left4_stem_1
30 activation_7
31 separable_conv_1_bn_reduction_right3_stem_1
32 separable_conv_1_bn_reduction_left4_stem_1
33 reduction_pad_1_stem_1
34 separable_conv_2_reduction_right2_stem_1
35 activation_9
36 activation_11
37 reduction_left2_stem_1
38 separable_conv_2_bn_reduction_right2_stem_1
39 separable_conv_2_reduction_right3_stem_1
40 separable_conv_2_reduction_left4_stem_1
```

In []:

```
In [87]: # we chose to train blocks, i.e. we will freeze  
# the first 249 layers and unfreeze the rest:  
for layer in model_NasNet.layers[:1031]:  
    layer.trainable = False  
for layer in model_NasNet.layers[1031:]:  
    layer.trainable = True
```

In []:

In []:

```
In [99]: # we need to recompile the model for these modifications to take effect  
# we use SGD with a low learning rate  
from keras.optimizers import SGD  
#model_NasNet.compile(optimizer=SGD(lr=0.0001, momentum=0.9), loss='sparse_categorical_crossentropy'  
  
model_NasNet.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

In []:

In []:

In [100]: # we train our model again (this time fine-tuning the top 2 inception blocks

alongside the top Dense layers

```
history_NasNet_2 = model_NasNet.fit_generator(  
    train_data_gen,  
    steps_per_epoch=total_train // batch_size,  
    epochs=10,  
    callbacks=[callback, cp_callback],  
    validation_data=val_data_gen,  
    validation_steps=total_val // batch_size  
)
```

Epoch 1/10

13/14 [=====>...] - ETA: 37s - loss: 0.2201 - acc: 0.9231

Epoch 00001: saving model to NasNet/cp_2.ckpt

14/14 [=====] - 690s 49s/step - loss: 0.2185 - acc: 0.9224 - val_loss: 0.5897 - val_acc: 0.7891

Epoch 2/10

13/14 [=====>...] - ETA: 36s - loss: 0.1974 - acc: 0.9251

Epoch 00002: saving model to NasNet/cp_2.ckpt

14/14 [=====] - 667s 48s/step - loss: 0.1912 - acc: 0.9278 - val_loss: 0.5748 - val_acc: 0.7969

Epoch 3/10

13/14 [=====>...] - ETA: 37s - loss: 0.1347 - acc: 0.9587

Epoch 00003: saving model to NasNet/cp_2.ckpt

14/14 [=====] - 681s 49s/step - loss: 0.1348 - acc: 0.9582 - val_loss: 0.6424 - val_acc: 0.8047

Epoch 4/10

13/14 [=====>...] - ETA: 36s - loss: 0.0974 - acc: 0.9735

Epoch 00004: saving model to NasNet/cp_2.ckpt

14/14 [=====] - 669s 48s/step - loss: 0.0945 - acc: 0.9743 - val_loss: 0.6915 - val_acc: 0.7969

Epoch 5/10

13/14 [=====>...] - ETA: 39s - loss: 0.0812 - acc: 0.9754

Epoch 00005: saving model to NasNet/cp_2.ckpt

14/14 [=====] - 701s 50s/step - loss: 0.0798 - acc: 0.9760 - val_loss: 0.6901 - val_acc: 0.7969

Epoch 6/10

13/14 [=====>...] - ETA: 36s - loss: 0.0580 - acc: 0.9858

Epoch 00006: saving model to NasNet/cp_2.ckpt

14/14 [=====] - 665s 48s/step - loss: 0.0585 - acc: 0.9851 - val_loss: 0.7155 - val_acc: 0.7969

Epoch 7/10

13/14 [=====>...] - ETA: 36s - loss: 0.0840 - acc: 0.9806

```
Epoch 00007: saving model to NasNet/cp_2.ckpt
14/14 [=====] - 667s 48s/step - loss: 0.0805 - acc: 0.9815 - val_loss: 0.7
181 - val_acc: 0.8086
Epoch 8/10
13/14 [=====>...] - ETA: 36s - loss: 0.0983 - acc: 0.9742
Epoch 00008: saving model to NasNet/cp_2.ckpt
14/14 [=====] - 665s 47s/step - loss: 0.0974 - acc: 0.9720 - val_loss: 0.7
794 - val_acc: 0.7969
Epoch 9/10
13/14 [=====>...] - ETA: 36s - loss: 0.0561 - acc: 0.9851
Epoch 00009: saving model to NasNet/cp_2.ckpt
14/14 [=====] - 659s 47s/step - loss: 0.0579 - acc: 0.9851 - val_loss: 0.7
514 - val_acc: 0.7969
Epoch 10/10
13/14 [=====>...] - ETA: 36s - loss: 0.0357 - acc: 0.9942
Epoch 00010: saving model to NasNet/cp_2.ckpt
14/14 [=====] - 659s 47s/step - loss: 0.0341 - acc: 0.9946 - val_loss: 0.9
113 - val_acc: 0.8086
```

In []:

In [101]: `test_loss, test_acc = model_NasNet.evaluate_generator(test_data_gen, verbose=0)`In [102]: `print('\nTest accuracy:', test_acc)`

Test accuracy: 0.8153078

In []:

In []: `#test_loss, test_acc = model_NasNet.evaluate_generator(test_data_gen, verbose=0)`In [90]: `#print('\nTest accuracy:', test_acc)`

Test accuracy: 0.8421052631578947

In []:

In []: